

Data & Things

(Spring 25)

Monday February 17

Lecture 8: Classification II

Jens Ulrik Hansen

Outline of this lecture

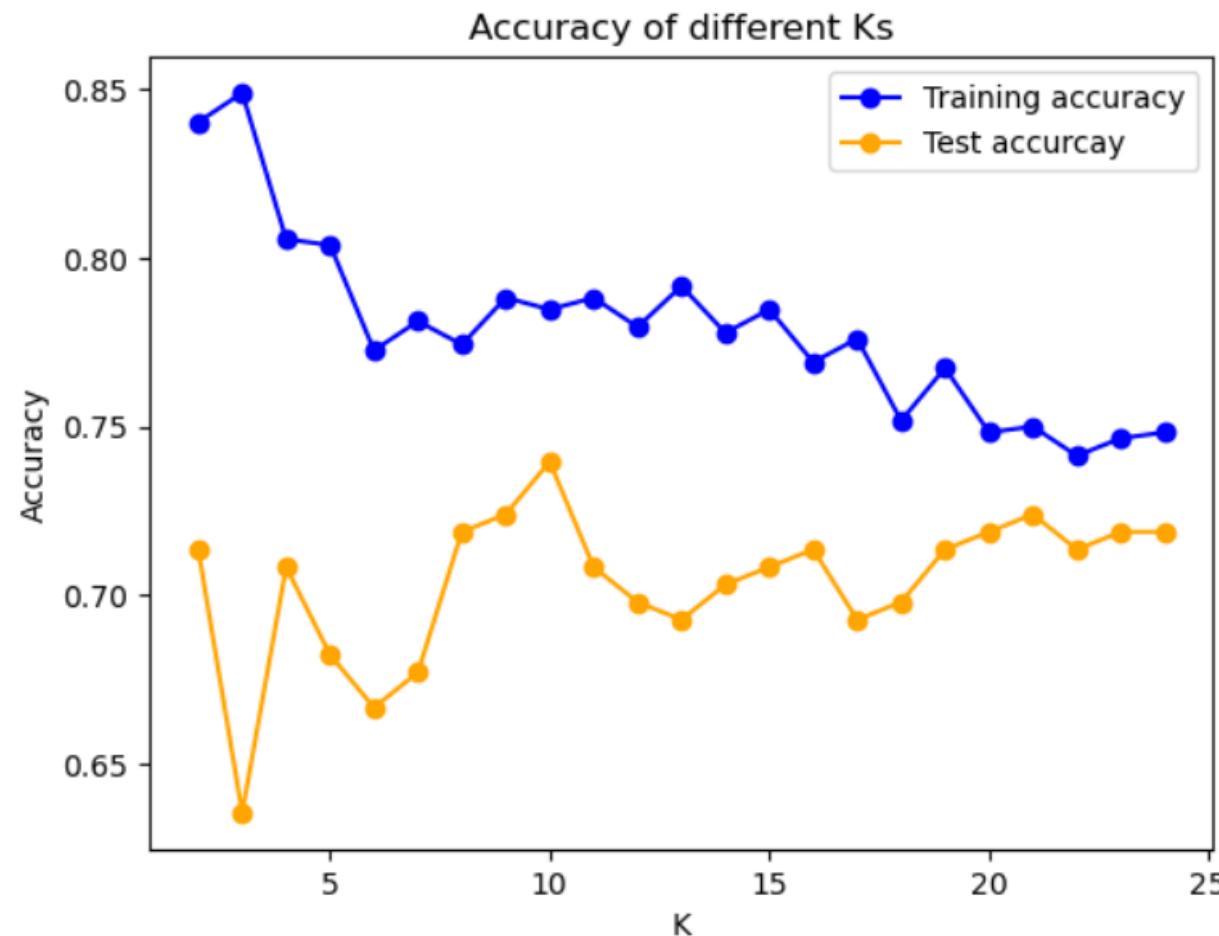
- Cross-validation
- Evaluating Classification models revisited
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting

Cross-validation

- Our train-test split from last time, is often not enough!
 - We might want train multiple models and pick the best one, or iteratively adjust “hyper-parameters” of our models
 - if use the performance on the test set to compare different models, we might **risk overfitting to the test set**
 - Moreover, the performance measure on the test set is no longer a good estimate of how the model will perform on new data
 - If we have smaller datasets, or use many variables in our model, we are left with a dilemma
 - We need a large portion of the data for training, to be able to get a good model
 - We need a large portion of the data for testing, to get a good estimate of how we will perform on new unseen data
 - Finally, train-test split can be quite sensitive to the actual random split (as we saw for the KNN example!)

Cross-validation

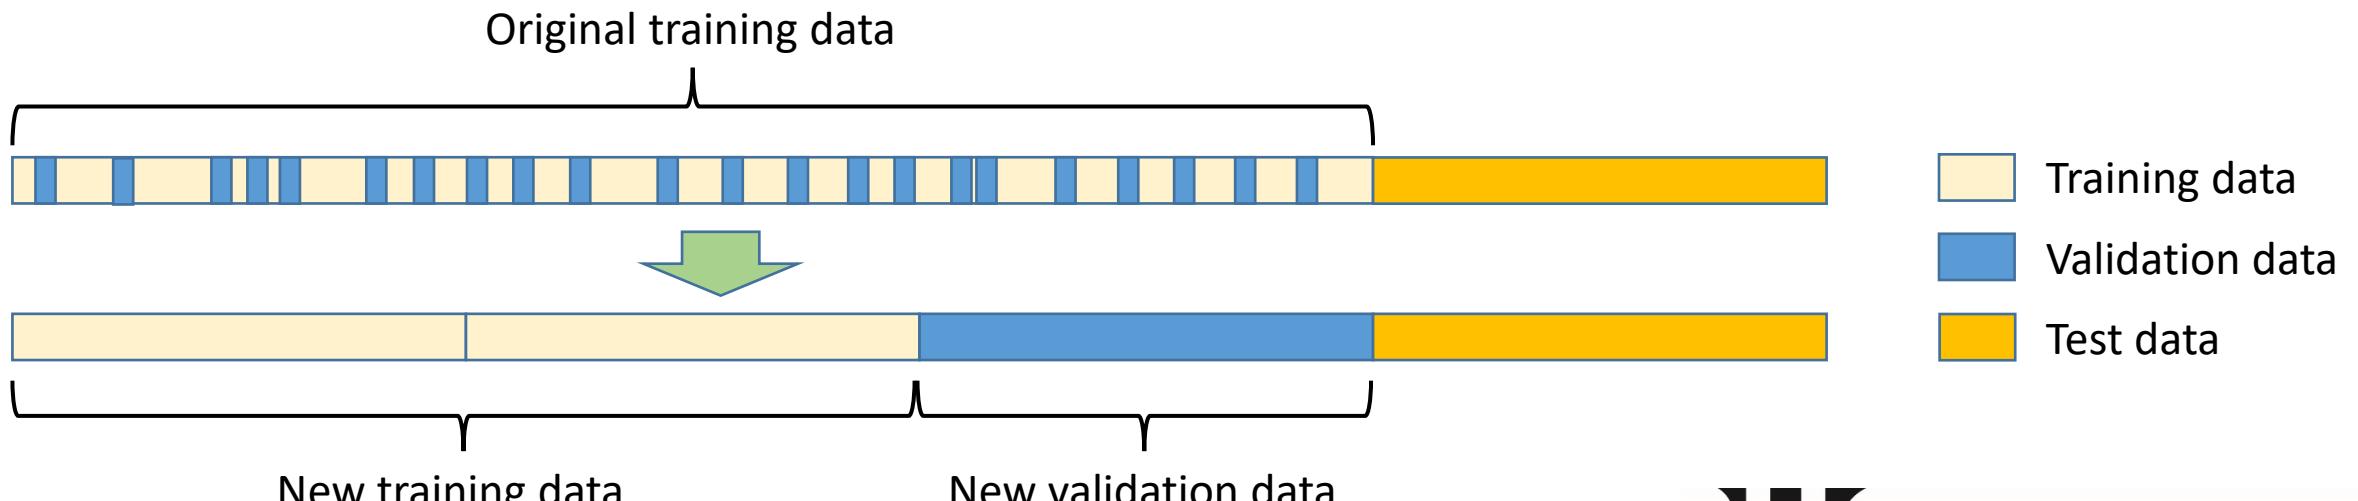
- Recall our KNN example on the diabetes example from last time...
 - We wanted to use the test accuracy to pick the right K
 - We had only 768 data points in total, so got few data points both for training and testing no matter how we did the split
 - Can we somehow utilize all the data for training and still get a good (unbiased) estimate of how well we will perform on future unseen data?



Cross-validation

- **Train-Validation-Test split**

- Use an additional **validation set** to evaluate and compare models, trying different models and setting to chose the best
 - More specific, we leave the test set as is and split the training set further into a new training set and a validation set.
 - We then train our models on the new training set
 - Evaluate, compare and chose final model based on performance on the new validation set
 - Finally, calculate an unbiased estimate of our model performance on new unseen data by calculating the performance on the test set
 - There is no one right proportion of splits, but an often sensible split is: 60% for training, 20% for validation, 20% for test

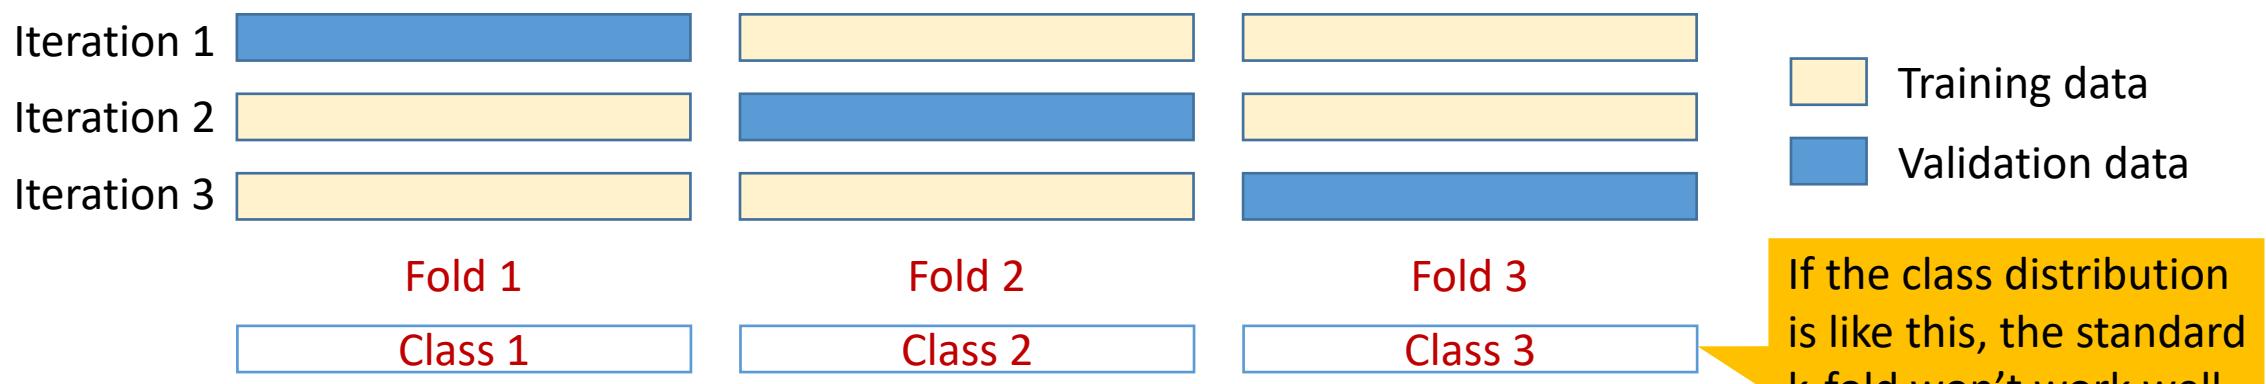


Cross-validation

- ***k*-fold cross-validation**

- ($k = 5$ or $k = 10$ are popular choices)
- Split the (train and validation) data D into k mutually exclusive subsets, each of approximately equal size: $D_1 \dots D_k$. Each D_i is called a *fold*.
- Do model construction and evaluation k times. Use the *average* accuracy.
 - At the i -th iteration, use fold D_i as the validation set and the others as the training set.

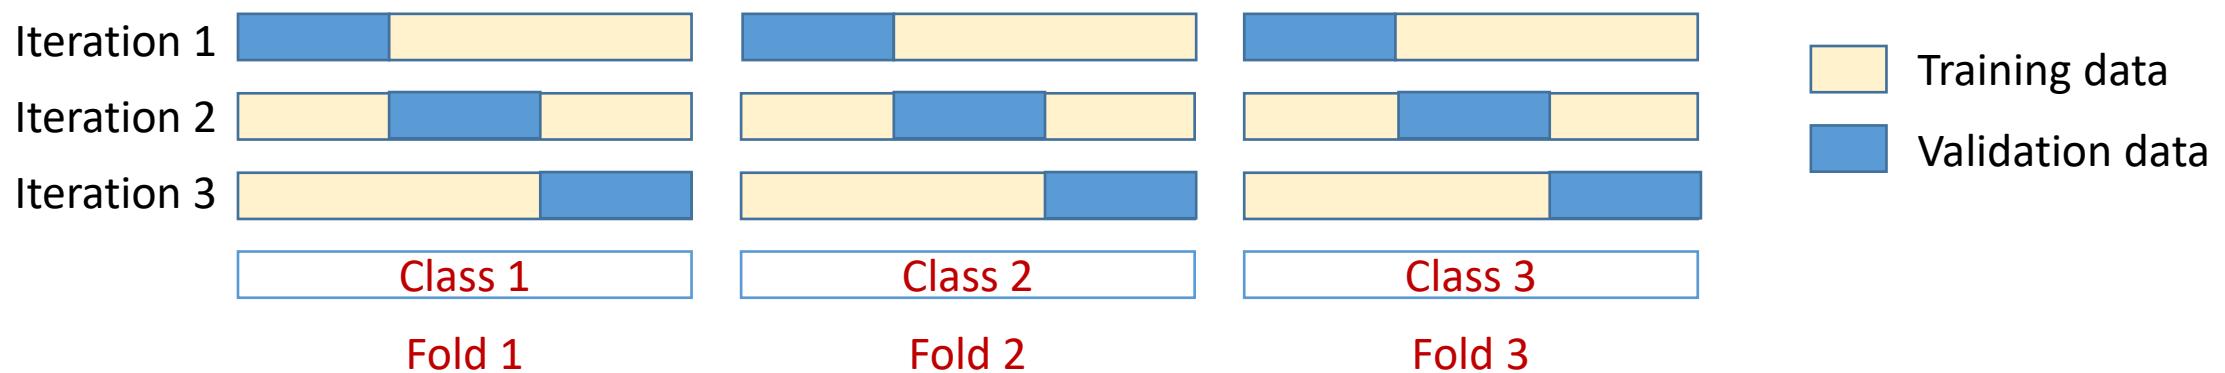
- Example of standard 3-fold cross validation



Cross-validation

- **Stratified cross-validation**

- folds are stratified so that *class distribution* in each fold is approximately the same as that in the initial given data.



- **Leave-one-out cross-validation**

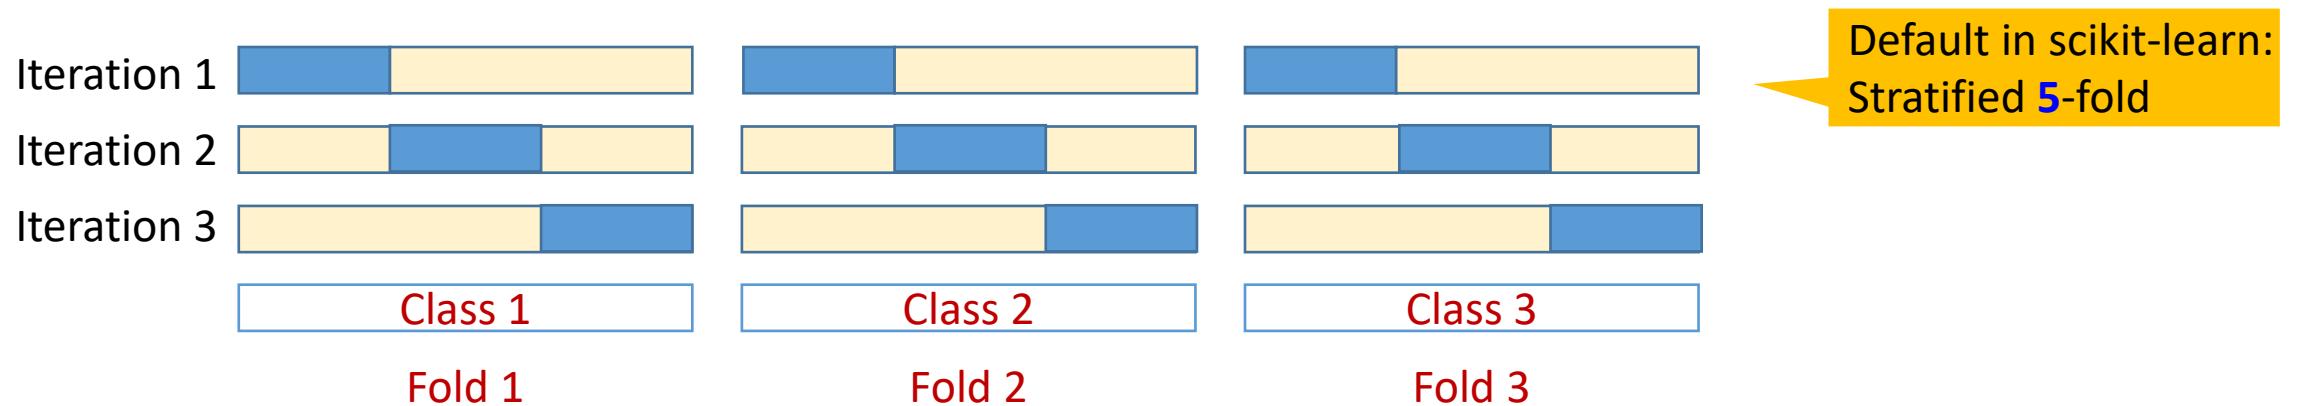
- k is set to the number of data points/rows in the dataset
- Use it only for small sized data; otherwise, too many models to construct!
- Often 5- or 10-fold cross validation is just as good or even better

Cross-validation

- Standard 3-fold cross validation



- Stratified 3-fold cross validation

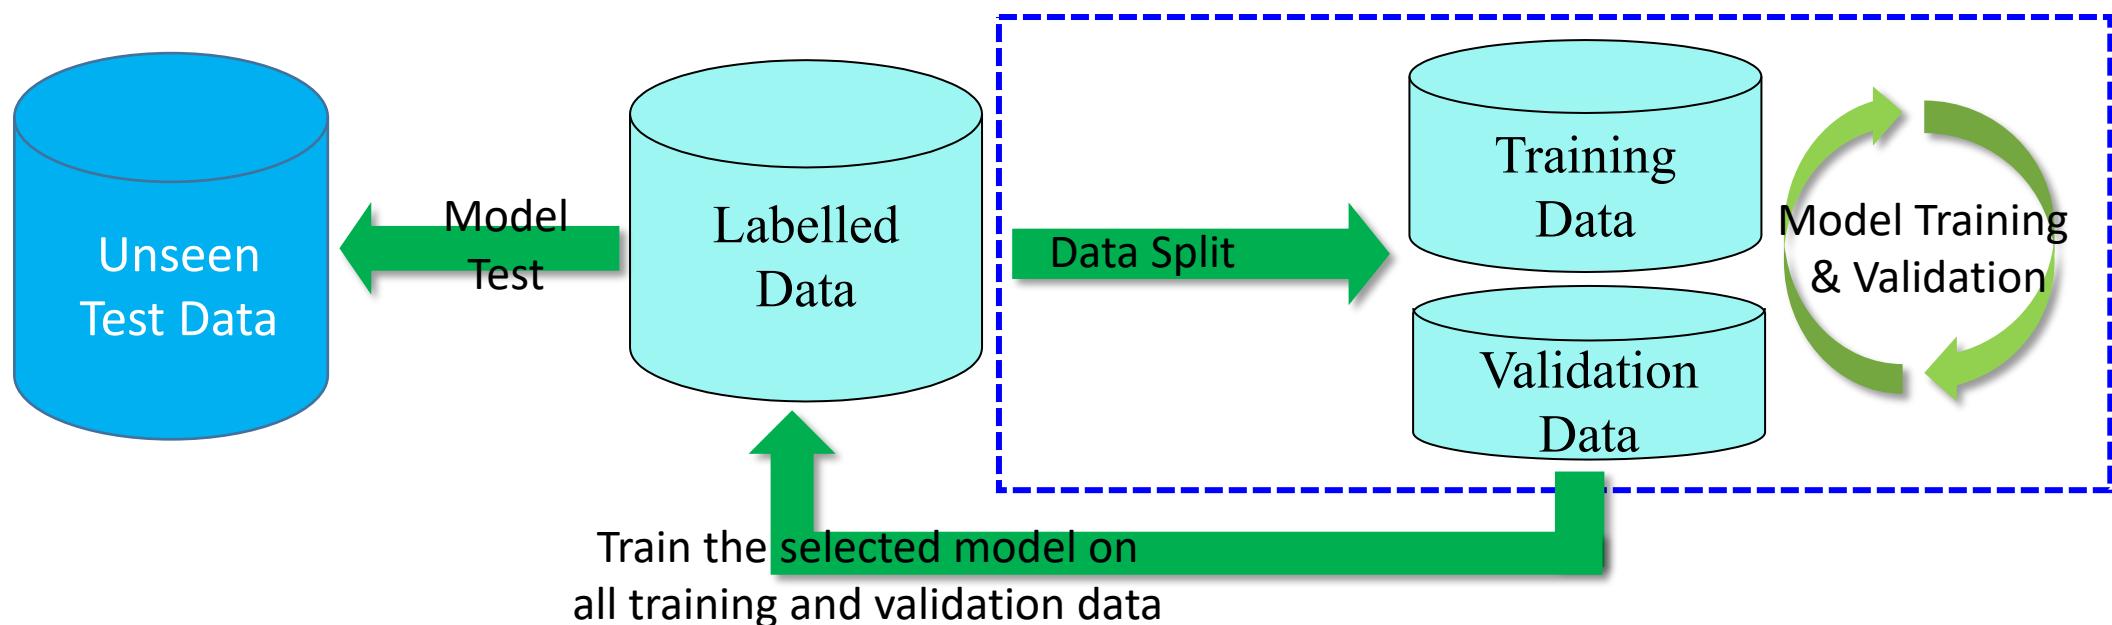


Cross-validation

- Cross-validation is not a way to construct an applicable model.
 - The function `cross_val_score(.)` builds multiple models *internally*, but these models are not returned.
- The purpose of cross-validation is to evaluate how well a *type* of model will generalize when it is trained on a specific dataset.
 - Model type: decision tree, random forest, KNN, SVM, ...
- Moreover, by using cross-validation, we can decide what type of model to use, and tune hyper-parameters for constructing a model
 - **Hyper-parameters**: algorithm parameters that can be set by the user before training a model.
 - E.g., `K` for KNN, `test_size` and `random_state` for `train_test_split(.)`, `gini` or `entropy` for a DT, ...
 - In contrast, **model parameters** are learned internally from training data
 - E.g., `a` and `b` coefficients in simple linear regression or logistic regression, weights in a neural network, ...

Cross-validation

- *After having used cross-validation to decide model type and tune hyper-parameters, we train the final model one last time on all the (training and validation) data (before evaluating it on the untouched test data)*

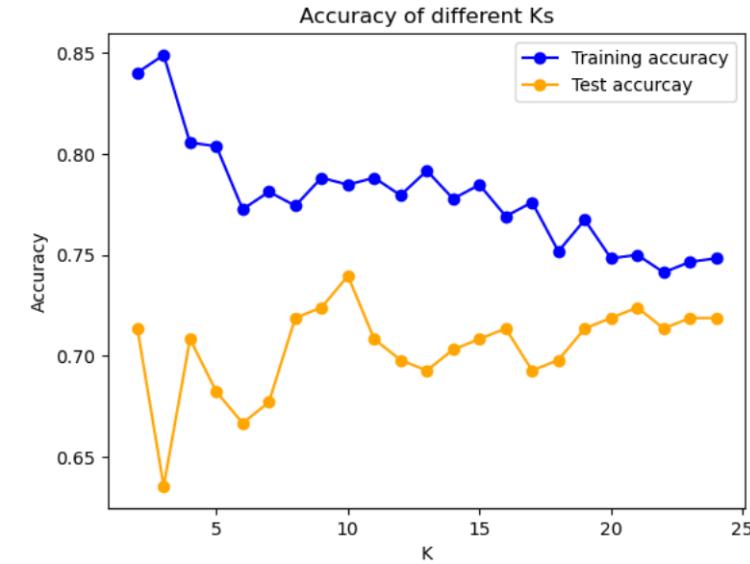


Outline of this lecture

- Cross-validation
- Evaluating Classification models revisited
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting

Evaluating Classification models revisited

- For classification models, we evaluated them by:
 - Accuracy: the fraction of correct predictions
 - Error rate: the fraction of incorrect predictions
- However, not all errors are equally important:
 - In medical diagnostics (1 = having the disease, 0 = not having it) two type of errors:
 - Predicted 1, but truly 0: Bad in the sense that it is not nice to be diagnosed with a disease, but later tests will likely be taken revealing that the patient did not have the disease after all.
 - Predicted 0, but truly 1: Much worse, the patient is led to believe she does not have the disease, and no further tests are made, which might lead to severe health conditions later.
 - In spam email detection (1 = it is spam, 0 = it is not spam) two types errors:
 - Predicted 1, but truly 0: Problematic if an actual mail is deleted as being spam.
 - Predicted 0, but truly 1: A little annoying to have a spam mail arrive in your inbox, but usually not a big problem



Evaluating Classification models revisited

- **The confusion matrix (binary case)**
 - Display predicted class versus actual class
 - **Accuracy:**
 - $(TP + TN) / (TP + FN + FP + TN)$
 - The fraction of the time the classifier makes the correct classification

	Predicted 1 (positive) (ex. spam)	Predicted 0 (negative) (ex. not spam)
Actual 1 (positive) (ex. spam)	TP (True Positive)	FN (False Negative)
Actual 0 (negative) (ex. not spam)	FP (False Positive)	TN (True Negative)

Evaluating Classification models revisited

- **The confusion matrix (binary case)**
 - Display predicted class versus actual class
 - Accuracy:
 - $(TP + TN) / (TP + FN + FP + TN)$
 - The fraction of the time the classifier makes the correct classification
 - **Precision (/positive predictive value):**
 - $TP / (TP + FP)$
 - The fraction of the positive classifications that is truly positive
 - If this is high, we have few real emails that are classified as spam

	Predicted 1 (positive) (ex. spam)	Predicted 0 (negative) (ex. not spam)
Actual 1 (positive) (ex. spam)	TP (True Positive)	FN (False Negative)
Actual 0 (negative) (ex. not spam)	FP (False Positive)	TN (True Negative)

Evaluating Classification models revisited

- **The confusion matrix (binary case)**
 - Display predicted class versus actual class
 - Accuracy:
 - $(TP + TN) / (TP + FN + FP + TN)$
 - The fraction of the time the classifier makes the correct classification
 - Precision (/positive predictive value):
 - $TP / (TP + FP)$
 - The fraction of the positive classifications that is truly positive
 - If this is high, we have few real emails that are classified as spam
 - Recall (Sensitivity/true positive rate):
 - $TP / (TP + FN)$
 - The fraction of the actual positive cases that the classifier detect as positive
 - If this is high, we have few cases of the disease that are not detected

	Predicted 1 (positive) (ex. spam)	Predicted 0 (negative) (ex. not spam)
Actual 1 (positive) (ex. spam)	TP (True Positive)	FN (False Negative)
Actual 0 (negative) (ex. not spam)	FP (False Positive)	TN (True Negative)

Evaluating Classification models revisited

- **The confusion matrix (binary case)**
 - Display predicted class versus actual class
 - Accuracy:
 - $(TP + TN) / (TP + FN + FP + TN)$
 - The fraction of the time the classifier makes the correct classification
 - Precision (/positive predictive value):
 - $TP / (TP + FP)$
 - The fraction of the positive classifications that is truly positive
 - If this is high, we have few real emails that are classified as spam
 - Recall (Sensitivity/true positive rate):
 - $TP / (TP + FN)$
 - The fraction of the actual positive cases that the classifier detect as positive
 - If this is high, we have few cases of the disease that are not detected
 - **F1 score:**
 - $(2 * \text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$

	Predicted 1 (positive) (ex. spam)	Predicted 0 (negative) (ex. not spam)
Actual 1 (positive) (ex. spam)	TP (True Positive)	FN (False Negative)
Actual 0 (negative) (ex. not spam)	FP (False Positive)	TN (True Negative)

Evaluating Classification models revisited

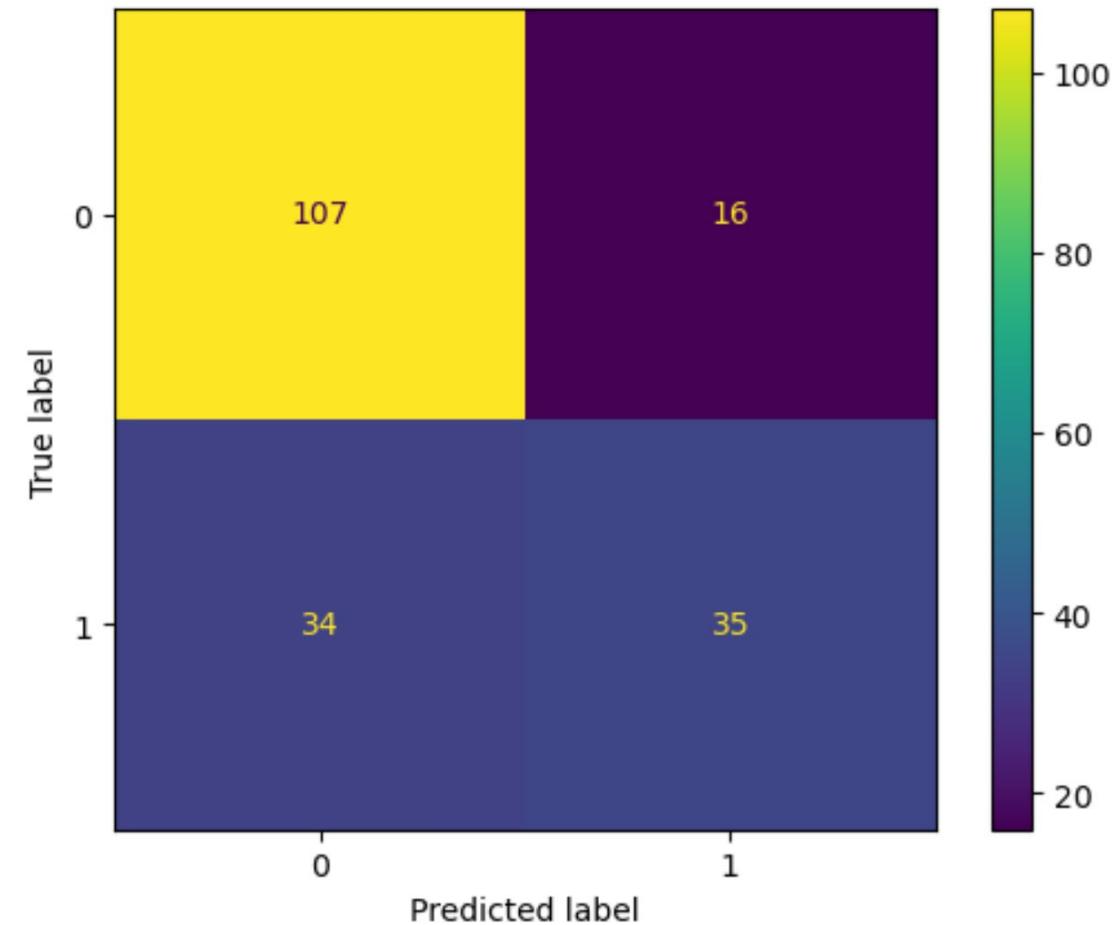
- **Other measures:**
 - **False positive rate:**
 - $FP / (FP + TN)$
 - **False negative rate:**
 - $FN / (FN + TP)$
 - **Specificity (/true negative rate):**
 - $TN / (TN + FP)$
 - **Negative predictive value:**
 - $TN / (TN + FN)$
 -

	Predicted 1 (positive) (ex. spam)	Predicted 0 (negative) (ex. not spam)
Actual 1 (positive) (ex. spam)	TP (True Positive)	FN (False Negative)
Actual 0 (negative) (ex. not spam)	FP (False Positive)	TN (True Negative)

Evaluating Classification models revisited

- Example: The KNN model (K=5) for the diabetes dataset from last time

K	Accuracy	Precision	Recall	F1
5	0.682292	0.557143	0.565217	0.561151



Evaluating Classification models revisited

- **Beyond the binary Confusion Matrix**

- Three pre-defined classes A, B, C
- Ground truth and classification result

Object	Ground Truth	Classification Result
object-1	A	A
object-2	B	A
object-3	C	C
object-4	C	C
object-5	B	B
object-6	A	B
object-7	B	B

- Confusion Matrix ($N \times N$ for N classes)

- The cell $M[i, j]$ counts the case that groundtruth i is classified as j

		Classification result			
		A	B	C	Total
Ground truth	A	1	1	0	2
	B	1	2	0	3
C	0	0	2	2	2
	Total	2	3	2	7

FP for A ←

FN for A →

TP for one class,
and TN for others

- **Accuracy** = $(TP+TN) / All$

- In this example, Accuracy = $(1+2+2)/7 = 71.4\%$

Evaluating Classification models revisited

- **Precision and Recall in general**

- Calculate the precision p_i for each class C_i
 - Overall precision is the *average* of all p_i 's
- Calculate the recall r_i for each class C_i
 - Overall recall is the *average* of all r_i 's
- Example:
 - $p_A = 30/60 = 1/2, r_A = 30/100 = 3/10$
 - $p_B = 60/120 = 1/2, r_B = 60/100 = 3/5$
 - $p_C = 80/120 = 2/3, r_C = 80/100 = 4/5$
 - Overall precision = 5/9,
overall recall = 17/30

Confusion matrix

Classification result

		A	B	C	Total
Ground truth	A	30	50	20	100
	B	20	60	20	100
C	10	10	80	100	
Total	60	120	120	300	

Recall
for A: r_A

Precision for A: p_A

Evaluating Classification models revisited

- **Analyze Your Confusion Matrix**

- Essentially, the more zeroes or smaller the numbers on all cells but the diagonal, the better a classifier is. So, you may analyze your confusion matrix and tweak your features accordingly.
- Confusion matrix gives strong clues as to where your classifier is going wrong.
 - E.g., if for Class A you can see that the classifier incorrectly predicts Class B for majority of the mislabeled cases, it indicates the classifier is somehow confused between classes A and B.
 - One way to fix this is to add discriminating features to improve classification of class A, e.g., more training data of A.

Evaluating Classification models revisited

- **Class imbalance**

- Rare positive examples but numerous negative ones, e.g., medical tests, fraud detection, etc. (...or the other way around)
- Traditional methods assume a balanced distribution of classes and equal error costs: not suitable for class-imbalanced data
- Typical methods for imbalance data in binary class classification:
 - **Oversampling**: re-sampling of data from positive class
 - **Under-sampling**: randomly eliminate tuples from negative class
 - **Threshold-moving**: moves the decision threshold, t , so that the rare class tuples are easier to classify, and hence, less chance of costly false negative errors
 - ...
- Still difficult for class imbalance problem on multiclass tasks

Evaluating Classification models revisited

- **Classification with class probabilities**

- All the evaluation metrics we have seen so far are evaluating classification models at set decision threshold (-we have set it)
- However, as we saw for logistic regression last time, some classification algorithms give us class probabilities instead of hard class labels
- But how do we choose the right threshold?
- Can we evaluate a model irrespective of what threshold we chose?

Evaluating Classification models revisited

- **Prediction probabilities** (like in logistic regression).
 - A probability of class membership is assigned instead of a label
 - A **threshold** can be used to control how to decide the predicted class label.

- Different thresholds lead to different metric values.
- This requires us to generate different confusion matrixes ☹

ID	Actual	Prediction Probability	>0.6	>0.7	>0.8	Metric
1	0	0.98	1	1	1	
2	1	0.67	1	0	0	
3	1	0.58	0	0	0	
4	0	0.78	1	1	0	
5	1	0.85	1	1	1	
6	0	0.86	1	1	1	
7	0	0.79	1	1	0	
8	0	0.89	1	1	1	
9	1	0.82	1	1	1	
10	0	0.86	1	1	1	

For positive label

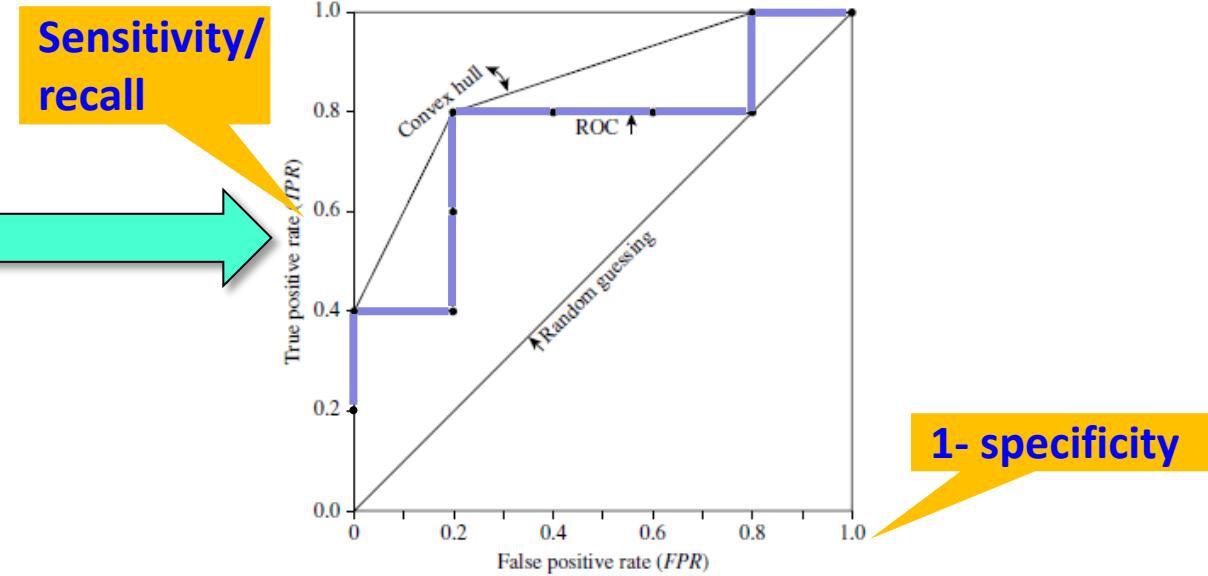
0.75	0.5	0.5	TPR
1	1	0.66	FPR
0	0	0.33	TNR
0.25	0.5	0.5	FNR

Evaluating Classification models revisited

- **Receiver Operating Characteristics (ROC)** curves: for visual comparison of binary classifiers
 - Rank your classification results in *descending* order of prediction probabilities
 - Calculate TPR and FPR for each current tuple in the ranked order
 - Mark each (FPR, TPR) point on the graph.
 - Connect all such points using a *convex hull*
- **NB:** TP, FP, TN, FN (and **TPR** and **FPR**) change as you seen more tuples in classification result

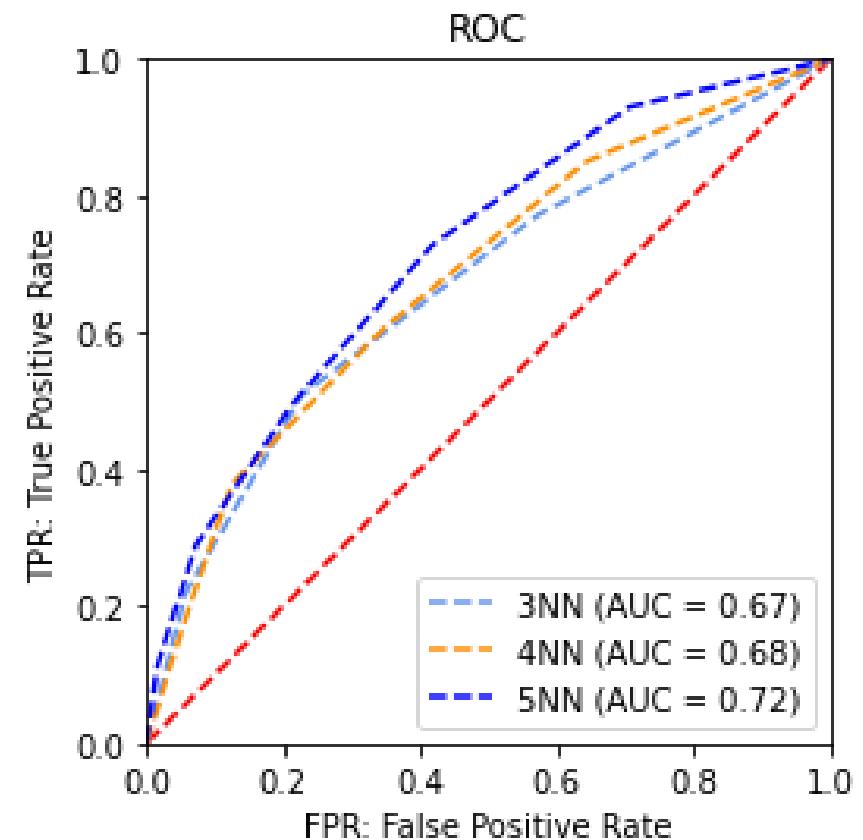


Tuple #	Class	Prob.	TP	FP	TN	FN	TPR	FPR
1	P	0.90	1	0	5	4	0.2	0
2	P	0.80	2	0	5	3	0.4	0
3	N	0.70	2	1	4	3	0.4	0.2
4	P	0.60	3	1	4	2	0.6	0.2
5	P	0.55	4	1	4	1	0.8	0.2
6	N	0.54	4	2	3	1	0.8	0.4
7	N	0.53	4	3	2	1	0.8	0.6
8	N	0.51	4	4	1	1	0.8	0.8
9	P	0.50	5	4	0	1	1.0	0.8
10	N	0.40	5	5	0	0	1.0	1.0



ROC Curves and AUC

- ROC Curves and AUC
 - A **ROC curve** shows the trade-off between the **True Positive Rate** and the **False Positive Rate**
 - The diagonal represents *random guessing*
 - The *area under the ROC curve* (**AUC**) is a measure of the “accuracy” of the model
 - The closer to the diagonal line (i.e., the closer the area is to 0.5), the less accurate is the model
 - A model with perfect accuracy will have an area of 1.0



Evaluating Classification models revisited

- **Examples**

- Let us look at the notebook “Evaluating classification models.ipynb”.

- **Exercise**

- Do Exercise 1 in the notebook “Exercises in Classification II.ipynb”

Outline of this lecture

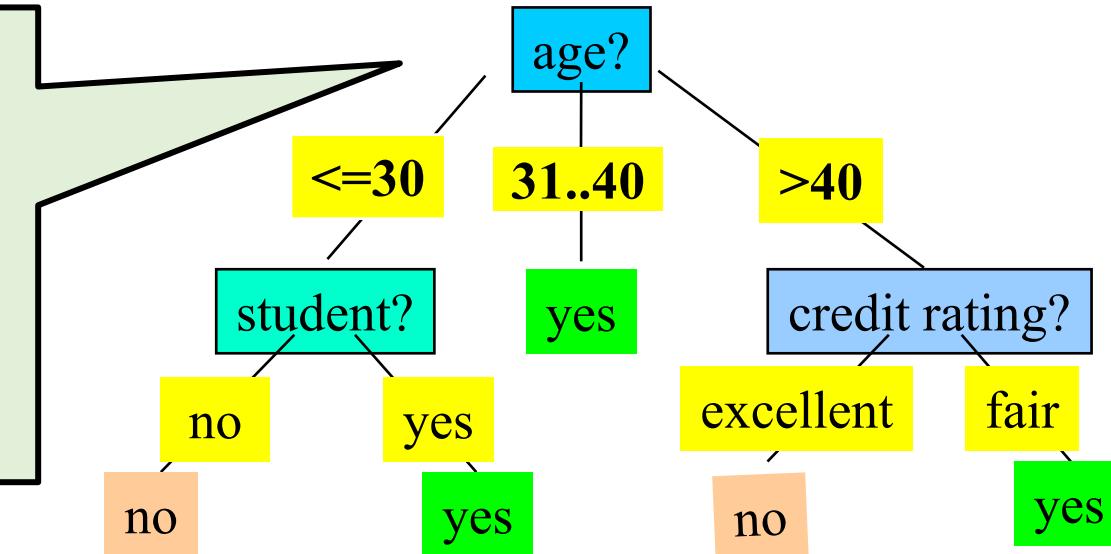
- Cross-validation
- Evaluating Classification models revisited
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting

Decision trees for classification

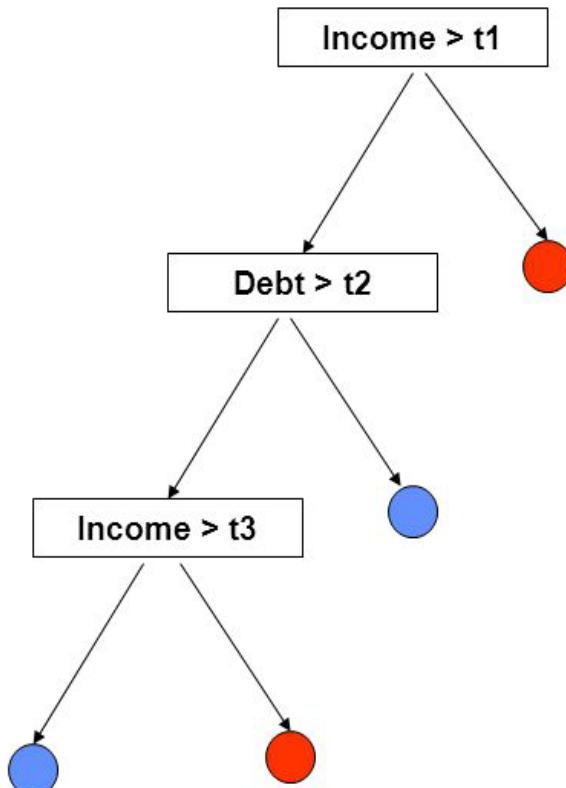
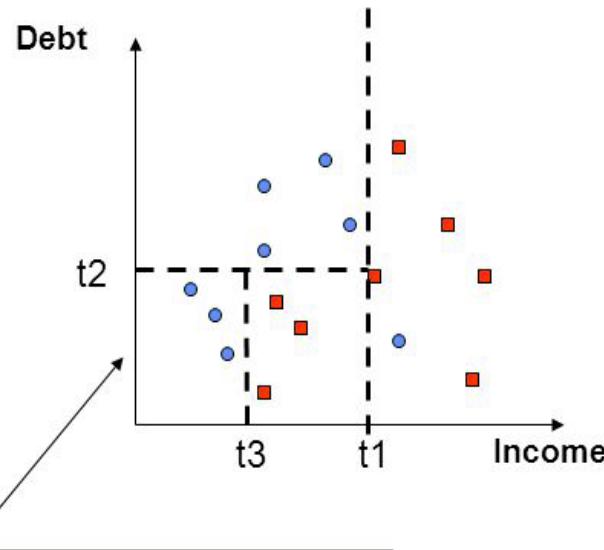
- **Decision Trees**

- This classification model is organized as a tree for decision making – it is thus called a **decision tree**.
- Internal nodes are associated with an *attribute/column* and arcs with *values* for that attribute.
- A leaf node tells the predicted class label (where the branches ends).

- Each person has attributes/columns:
 - (age, student [yes/no], credit rating)
 - p1(18, yes, fair)
 - p2(55, no, excellent)
- **Two classes**
 - **Buy computer**
 - Not buy computer



Decision trees for classification



- **Training a decision tree model**

- Recursive binary splits top-down.
- In a 2D feature dataset, each split corresponds to drawing a horizontal or vertical line.
- Splits are chosen to minimize either *Gini index* or *Entropy* (both measures of node impurity).
- Each final region corresponds to a leaf in the decision tree.
- For each region, the predicted class corresponds to the most prevalent class – class probabilities can be obtained by noting the fraction of each class.

<https://github.com/martian1231/decisionTreeFromScratch>

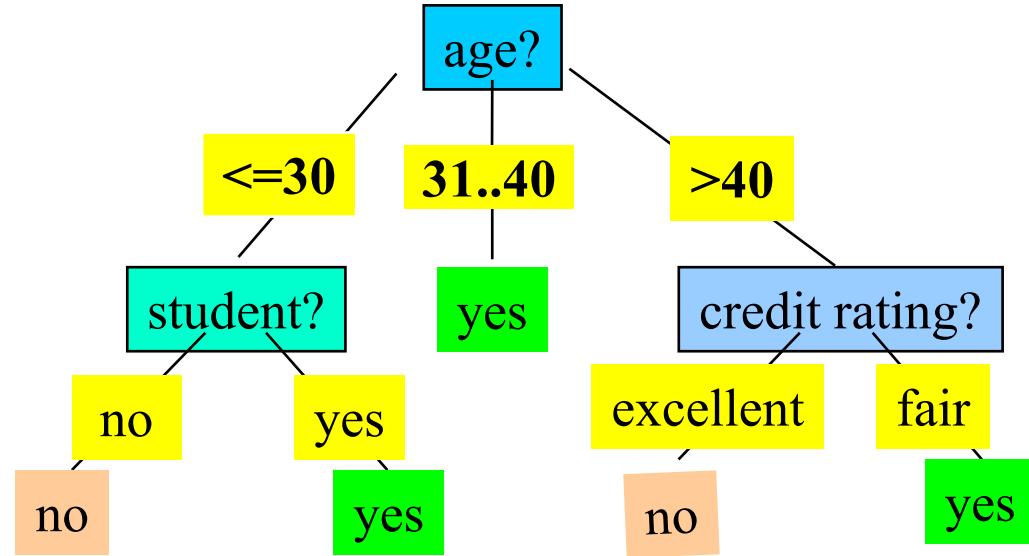
Decision trees for classification

- **Pros and Cons of decision trees**

- Pros: Easy to interpret (even by non experts) and visualize
- Pros: Does not require much preprocessing such as scaling or creation of dummy variables for categorical variables
- Pros: Highly flexible
- Cons: Can grow big and then becomes hard to interpret
- Cons: Does not have the same level of predictive accuracy as other models
- Cons: Have high variance

- **Decision trees for regression**

- The prediction is made by taking the average of values in a region
- Instead of minimizing Gini index or entropy, we minimize a local version of RSS.



Decision trees for classification

- **Examples**

- Let us look at the notebook “Decision trees and ensemble models.ipynb”.

Outline of this lecture

- Cross-validation
- Evaluating Classification models revisited
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting

Ensemble models using bagging and boosting

- ***Ensemble models*** combine many simple models into one single potential powerful model
- Combining decision trees for classification (or regression) turns out to good candidates for such simple models
- Popular ensemble methods:
 - ***Bagging***: averaging the prediction over a set of independent classifiers
 - ***Random Forest***: a random set of decision trees
 - ***Boosting***: weighted vote with a set of dependent classifiers

Ensemble models using bagging and boosting

- **Bagging: Bootstrap aggregating**

- Analogy: Diagnosis based on the *majority vote* of multiple doctors trained from samples of the same superset.
- Reduce variance of models with high variance, such as decision trees
- Training: Given a set D of d tuples, generate m new training data sets D_i 's
 - Each training set D_i of d' tuples is sampled with *replacement* from D (i.e., **bootstrap**)
 - Set D_i is used to train a classifier model M_i
- Classification: classify an unknown sample x using all m models (M_1 to M_m)
 - Each classifier M_i returns its class prediction
 - The bagged classifier M^* counts the votes and assigns the class with the most votes to x
- Accuracy
 - Often significantly better than a single classifier derived from D
 - For noisy data: not considerably worse, more robust

Ensemble models using bagging and boosting

- **Random Forest (of Decision Trees)**

- An improved version of bagging
 - In addition to bootstrap sampling, we are also sampling a subset of features for each decision tree
- Formally, for each decision tree
 - We sample with replacement a subset of the original training set
 - We sample set of m of the feature variables – if there is a total of p features, a common choice for m is \sqrt{p}
 - We then train the decision tree on this subset of the training dataset and only allows for split involving the sampled m features.
- The rationale behind Random Forest
 - That we also sample only a small subset of features for each decision tree makes the decision trees much more varied and uncorrelated, which in turn make the average more robust.
- Decision rule
 - To make a prediction, Random forest average over all predictions from all the trees for regression and take the most popular vote for a class among all the trees for classification

Ensemble models using bagging and boosting

- **Boosting**

- Analogy: Consult several doctors, based on a combination of weighted diagnoses—weight assigned based on the *previous* diagnosis accuracy
- How boosting works?
 - Weights are assigned to each training tuple
 - A series of k classifiers is iteratively learned
 - After a classifier M_i is learned, the weights are updated to allow the subsequent classifier, M_{i+1} , to pay more attention to the training tuples that were misclassified by M_i
 - The final M^* combines the votes of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Compared with bagging: Boosting tends to have greater accuracy, but it also risks overfitting the model to misclassified data
- Popular variations of boosting: ***AdaBoost*** and ***XGBoost***

Ensemble models using bagging and boosting

- **Interpretability of ensemble models**

- As we are combining many models to achieve higher accuracy, interpretability is no longer straight-forward – a classic example of the trade-off between predictive accuracy and interpretability
- However, for ensemble models based on decision trees, we can record in each decision and for each feature variable X_i the drop in Gini index (or local RSS) and average this across all the trees in the ensemble – this will give us the ***variable importance*** of the feature variable X_i .
 - In this way, we can get a ranking of which of the features are most important relative to each other in the predictions the ensemble model makes.

Summary on Decision trees and ensemble models

- Decision trees are easy to understand, easy to train, and easy to interpret
- Decision trees are flexible and easily handle both numeric and categorical variables (without scaling) and can be used for both regression and classification
- Decision trees often have limited accuracy and very high variance
- For these reasons, Decision trees are ideal as the simple models combined in ensemble models
- Random Forest is an ensemble method that often perform well and rarely overfit
- Boosting methods are more prone to overfitting, but can also achieve higher predictive accuracy than Random Forest

Ensemble Learning: Increasing the Accuracy

- **Examples**

- Let us look at the notebook “Decision trees and ensemble models.ipynb”.

- **Exercise**

- Do Exercise 2 in the notebook “Exercises in Classification II.ipynb” – the same as Exercise 2 in the notebook “Exercises in Classification I.ipynb” from last time