

# Data & Things

(Spring 25)

---

Friday February 21

**Lecture 10: Clustering**

Jens Ulrik Hansen

# Introduction to machine learning

---

- **Supervised learning** generalizes from *Labeled data* to facilitate future predictions of label based on input data.
  - Given a collection of feature variables  $X_1, X_2, \dots, X_p$ , can we find a model that approximately predict a response variable  $\hat{Y}$ ?
  - Two stages: Training and prediction
  - **Classification:** Predict a *discrete* value from a *pre-defined* set of class labels
  - **Regression:** Predict a *continuous* value from a continuous range
- **Unsupervised learning** find patterns in *unlabeled data*. It works on input data only.
  - *We are only given the feature variable  $X_1, X_2, \dots, X_p$  and no response variable or true labels  $Y$  – we have to look for patterns in  $X_1, X_2, \dots, X_p$  without having an example of what we are looking for and thereby no obvious way of testing model performance*
  - Only one stage, work on the entire dataset
  - E.g., clustering, customer segmentation, outlier detection for website access patterns, etc.

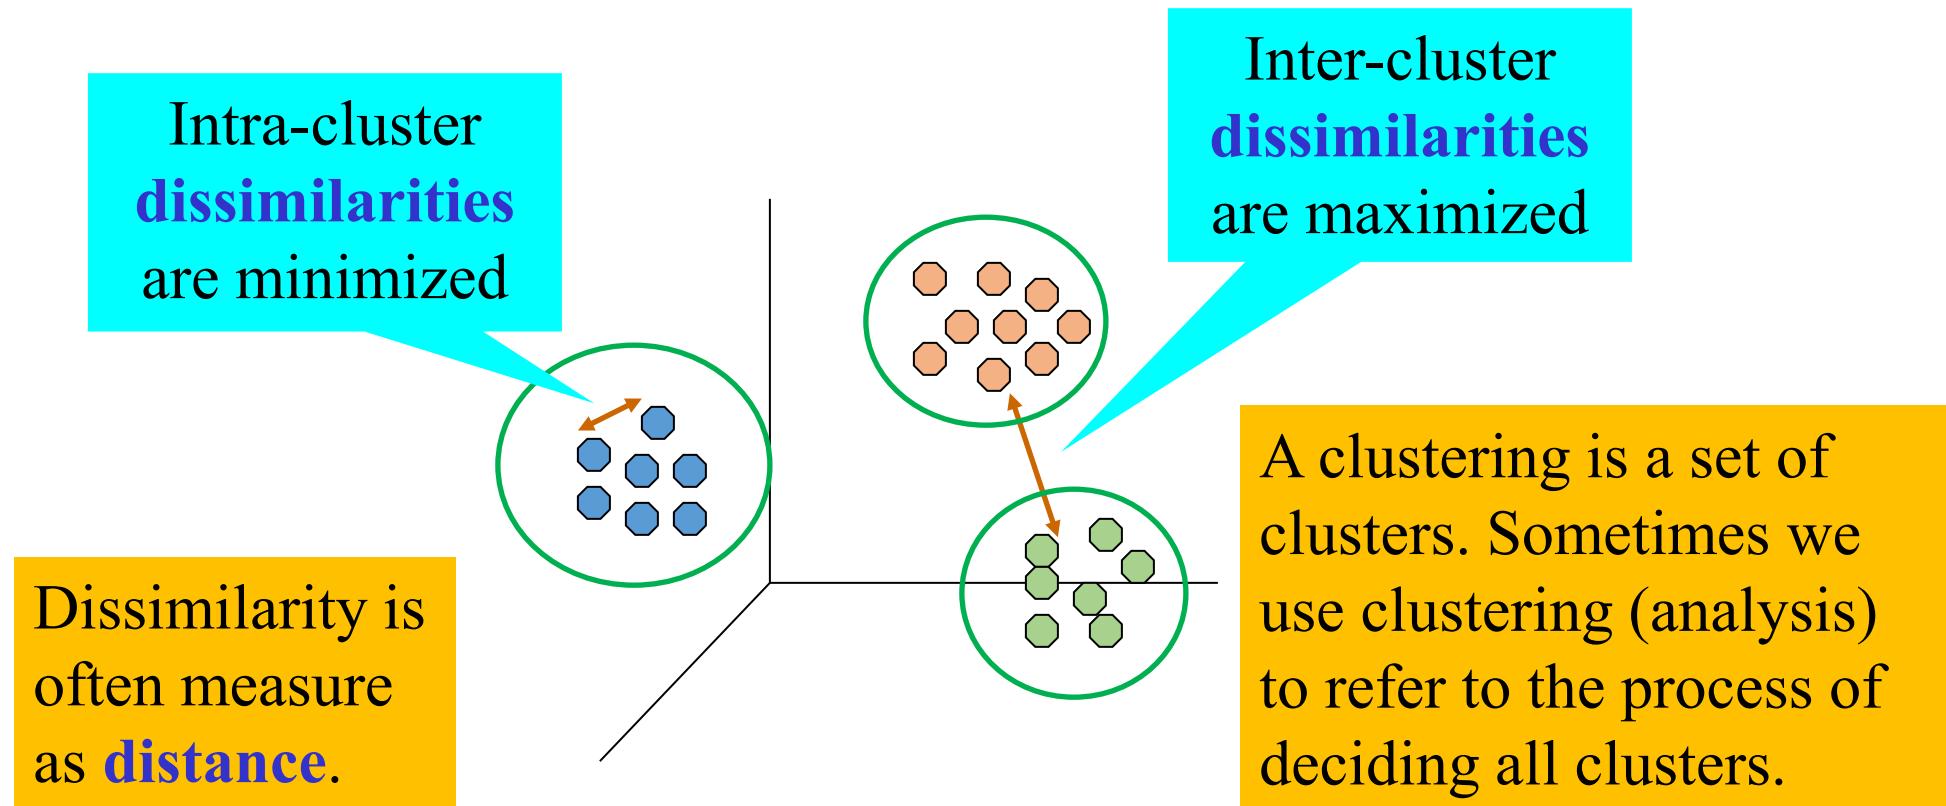
# Outline of this lecture

---

- Clustering as an example of unsupervised learning
- K-means clustering
- Hierarchical clustering
- DBSCAN clustering
- Evaluation of clustering models

# Clustering as an example of unsupervised learning

- **Clustering:** Grouping of objects, s.t. the points in a group (*cluster*) are similar (or related) to each other and different from (or unrelated to) points in other groups



# Clustering as an example of unsupervised learning

- Another example (for two feature variables):

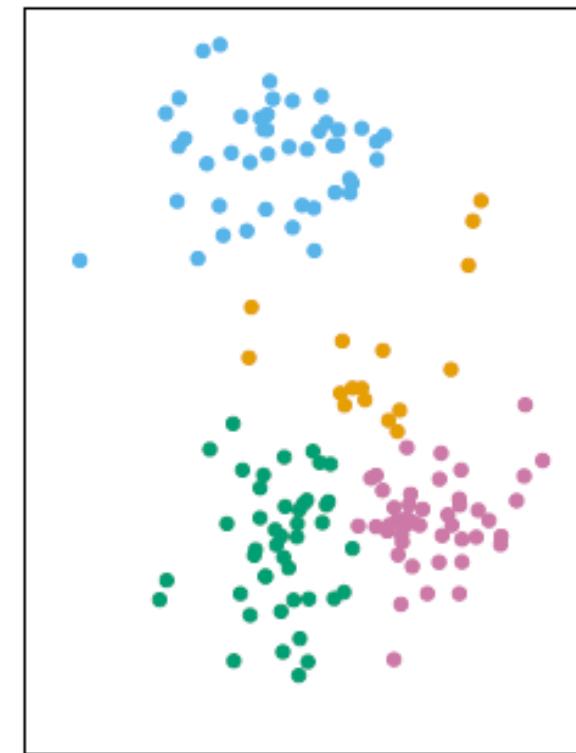
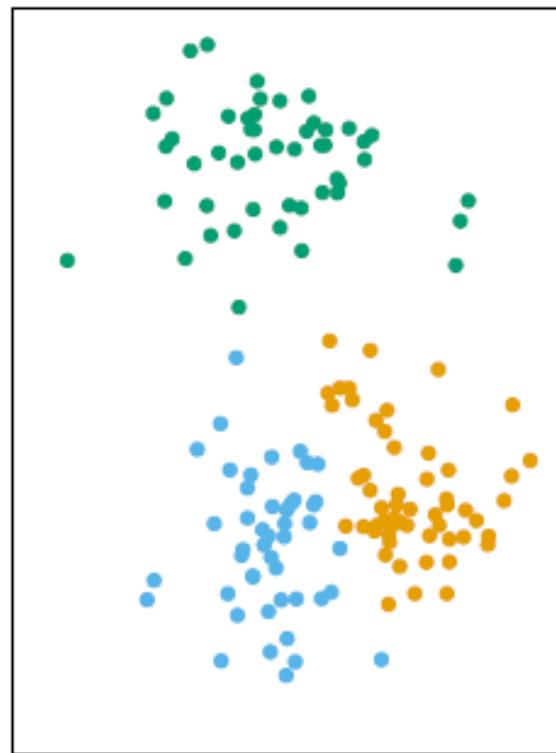
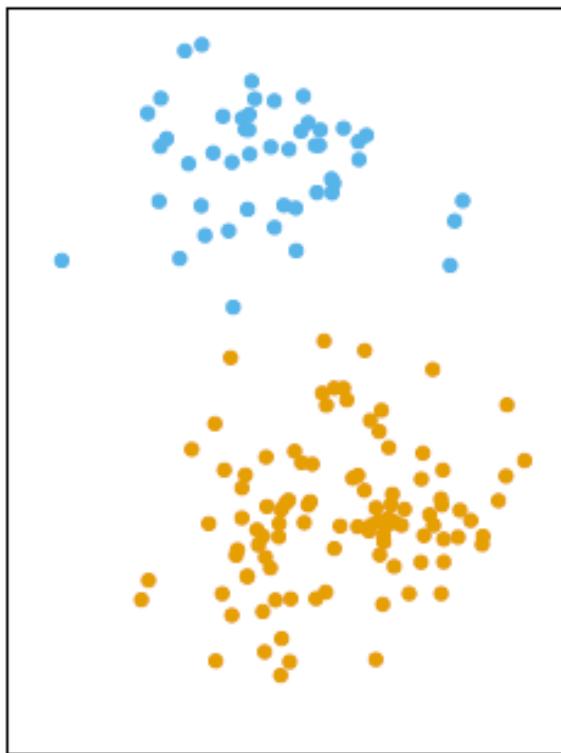


Figure 12.7 from James, G., Witten, D., Hastie, T., Tibshirani, R., and Taylor, J. (2023). *An Introduction to Statistical Learning – with Applications in Python*. Springer

# Clustering as an example of unsupervised learning

---

- **Formal Definition of Clustering**

- **Input:** A collection  $C$  of data objects
- **Output:** A set of *disjoint* clusters whose union is  $C$ .
  - Objects in the same clusters are *similar* to each other.
  - Objects in one cluster are *dissimilar* to those in other clusters.
- **Process:** Compute similarities between data points and group similar data points into the same cluster.
- Typical use of clustering
  - As a *stand-alone tool* to get insight into data distribution
  - As a *preprocessing step* for other algorithms
- ***Unsupervised learning: clusters are not pre-defined***

# Clustering as an example of unsupervised learning

---

## Classification (supervised)

- Predefined classes
  - Number of classes
  - Meaning of classes
- Work for any number of points (in the prediction stage)
  - Given a point, a classifier (trained model) assigns it to a class

## Clustering (unsupervised)

- No prior knowledge about
  - Number of clusters \*
  - Meaning of clusters
- There must be a sufficient number of points
  - Meaningless to conduct clustering analysis on one or few objects

# Clustering as an example of unsupervised learning

---

- **Examples of application of clustering**

- *Market segmentation* (Customer segmentation): If we can group our customers into groups of similar people, based on demographics or shopping behavior for instance, we can target the groups separately with different marketing campaigns.
- *Image segmentation*: We can use clustering to segment images into different regions, for instance useful in Medical Image Analysis.
- *Gene analysis*: Group similar gene expressions to understand genes' functions and regulatory mechanisms
- *Topic Modeling*: Group a collection of text document (like news articles) into different topics (non-predefined topics)
- ...

# Clustering as an example of unsupervised learning

- **Basic Steps of Clustering**

1. Feature selection
2. Proximity measure
  - Similarity of two feature vectors
3. Clustering criterion
  - Expressed via a cost function or some rules
4. Clustering algorithms
  - Choice of algorithms
5. Validation of the results
  - Validation test (also, *clustering tendency* test)
6. Interpretation of the results
  - Integration with applications

• What attributes should we consider?

• How to measure similarity?

• How close two points should be to get into the same cluster?

Domain expertise may be needed.

# Clustering as an example of unsupervised learning

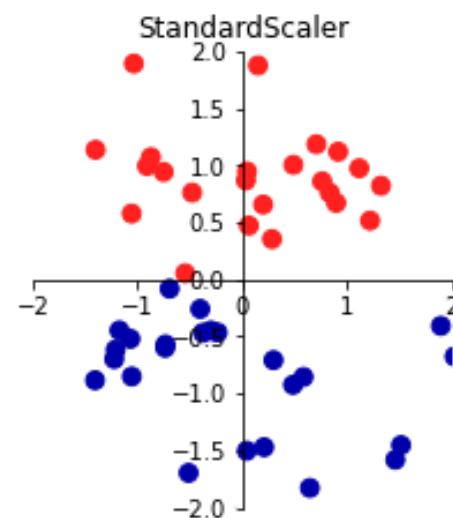
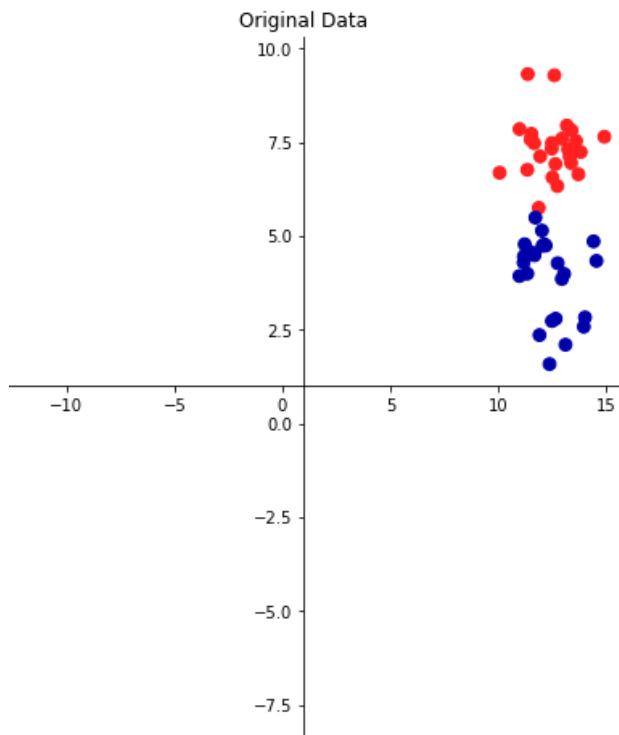
- **Data scaling before Clustering**

age	income
64	87083.24
33	76807.82
24	12043.60
33	61972.00
78	60120.32
62	40058.42

- If we calculate distance directly on this dataset, the distance will very likely be dominated by the income values.
  - Dimensions age and income are not measured in the same scale.
- Data (re)scaling is needed before reasonable distances can be calculated on the two dimensions.
  - This is part of preprocessing of the data before distance-based ML algorithms, e.g., kNN for classification and those for clustering

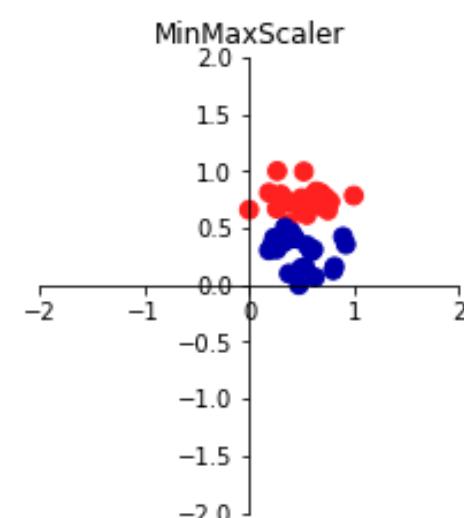
# Clustering as an example of unsupervised learning

- **Preprocessing and Scaling**



Standard Scaling (aka standardization or Z-score normalization)

- Afterwards, for each feature has **mean=0** and **variance=1**



Min-Max Scaling (aka Normalization)

- Shifts the data, *s.t.* each feature falls in **[0..1]**

# Outline of this lecture

---

- Clustering as an example of unsupervised learning
- K-means clustering
- Hierarchical clustering
- DBSCAN clustering
- Evaluation of clustering models

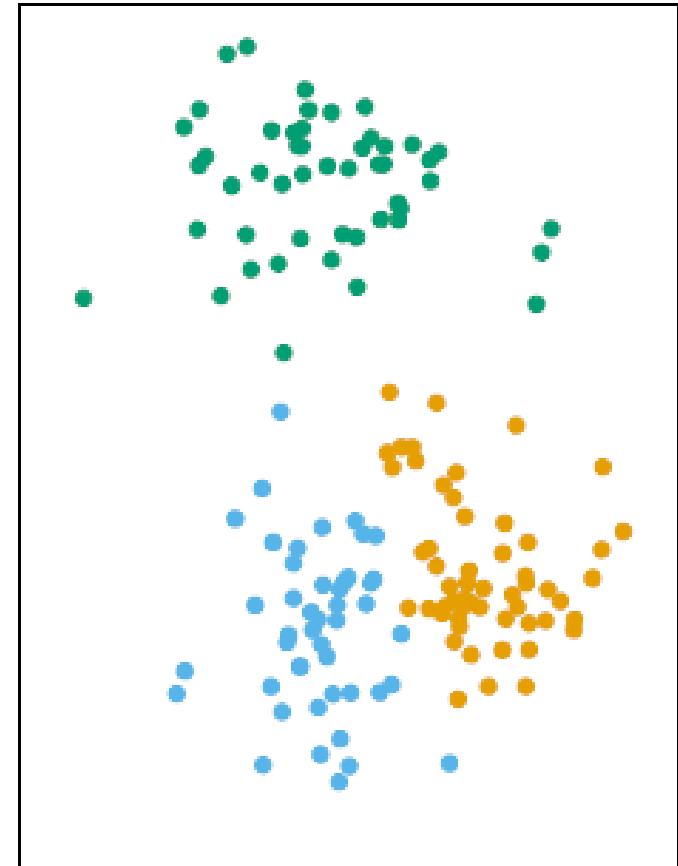
# K-means clustering

- A clustering is a set of clusters (set of points) such that:
  - All data points belong to exactly one cluster
- **The K-means clustering problem**
  - Can we choose K clusters such that we minimize the *within-cluster variation*?

$$\underset{C_1, \dots, C_K}{\text{minimize}} \left\{ \sum_{k=1}^K W(C_k) \right\}$$

- If we measure within-cluster variation with squared Euclidian distance, it reduces to:

$$\underset{C_1, \dots, C_K}{\text{minimize}} \left\{ \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2 \right\}$$



# K-means clustering

- Given K, the ***K-means algorithm*** works in four steps:

Initialization

- Partition all points *randomly* into K nonempty subsets (clusters)

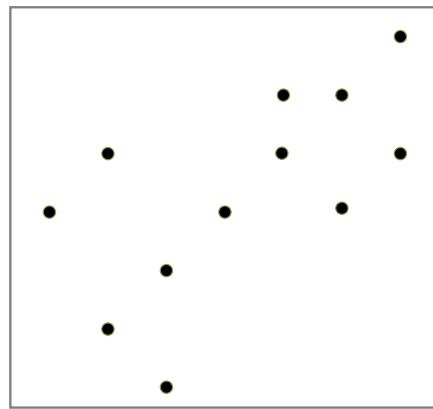
Iterations

- Compute the **centroids** of the clusters of the current partitioning
  - The centroid is the center, i.e., **mean point**, of a cluster
- Assign each point to the cluster with the *nearest* centroid
- Go back to Step 2, repeat and stop when the assignment does not change, or the change is sufficiently small
  - Convergence

Convergence

# K-means clustering

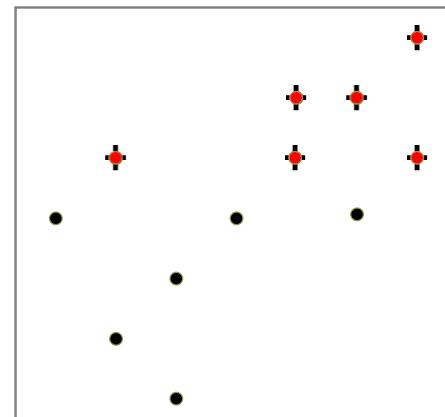
## An Example of K-Means Clustering



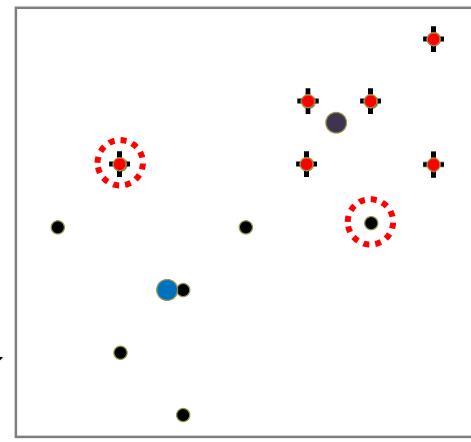
The initial data set

K=2

Arbitrarily partition points into K groups

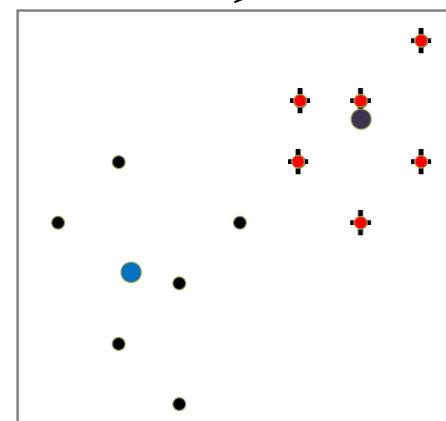


Update the cluster centroids

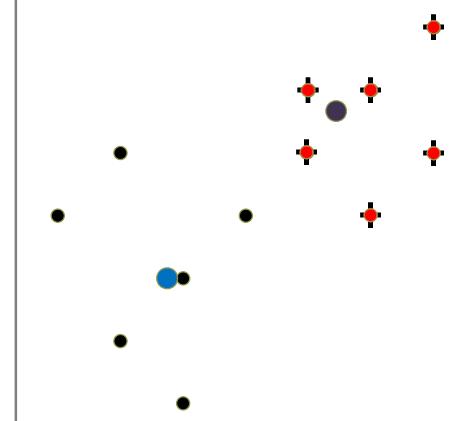


Loop if needed

Reassign points



Update the cluster centroids



Partition points into  $k$  nonempty subsets

**Repeat**

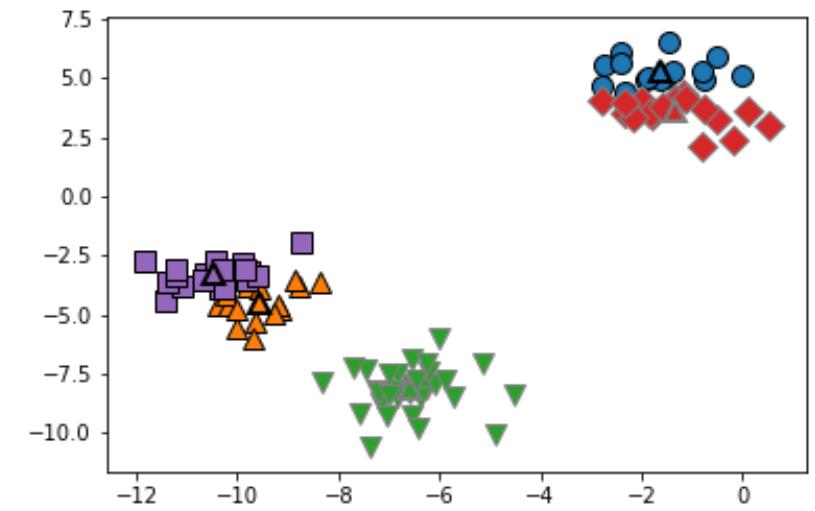
- Compute centroid (i.e., mean point) for each cluster
- Assign each point to the cluster of its nearest centroid

**Until** convergence

# K-means clustering

- **Notes on K**

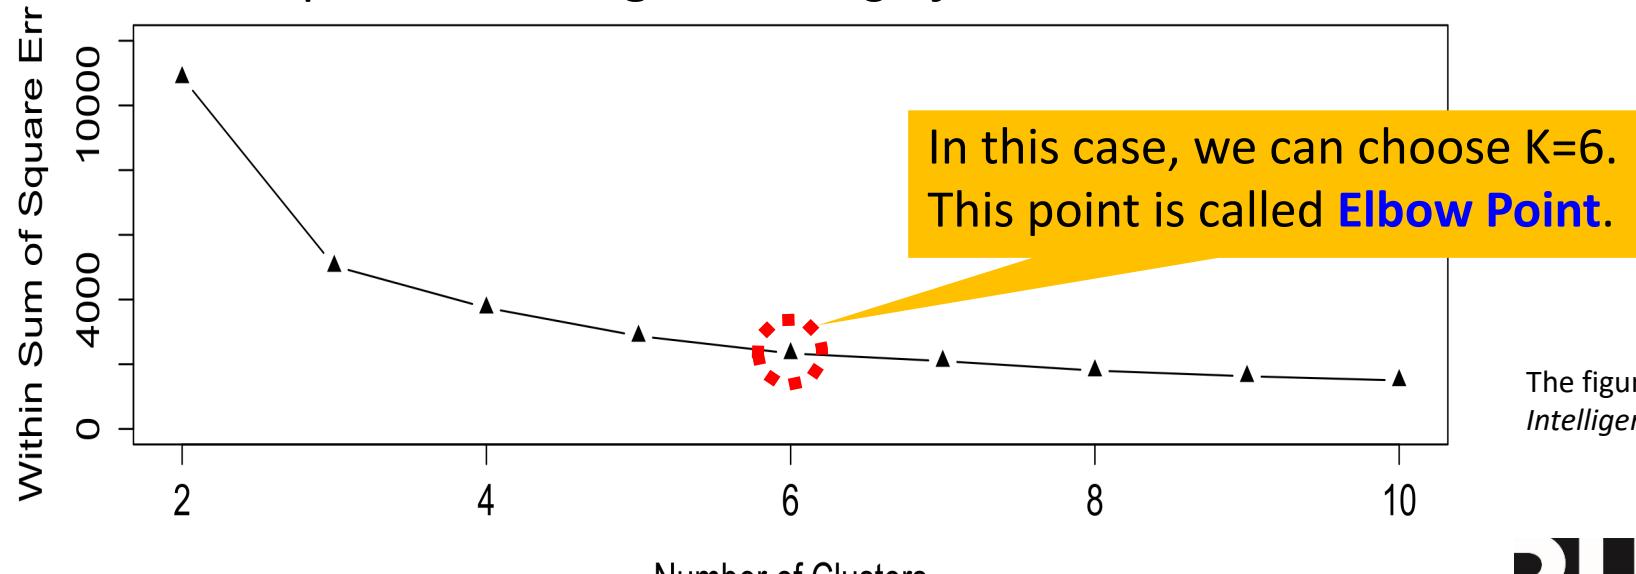
- Different initializations, might affect the final clustering – K-means is only locally optimal
- The time complexity of K-means depends on K
- A larger K:
  - More clusters to maintain, more mean points to calculate, and more distance calculations and comparisons in the reassignment step.
- A smaller K:
  - Less clusters to maintain, less mean points to calculate, and less distance calculations and comparisons in the reassignment step.
- K may also affect the clustering quality
- We may use EDA and visualization to decide K.
  - Only useful if the data is not high-dimensional



# K-means clustering

- **Elbow Method: To decide the best K**

- Let  $c_i$  be the *centroid/mean* of cluster  $C_i$  in a given clustering result.
- We check the **Sum of Squared Distance** (aka sum of squared error **SSE**) for all points  $p$  in all clusters:  $E = \sum_{i=1}^k \sum_{p \in C_i} (p - c_i)^2$
- Vary K from 1 to a max (e.g., 10), plot a graph for (K, SSE), and find the K value *after which the performance gain is insignificant*.



The figure is from *Introduction to R for Business Intelligence* by Jay Gendron

# K-means clustering

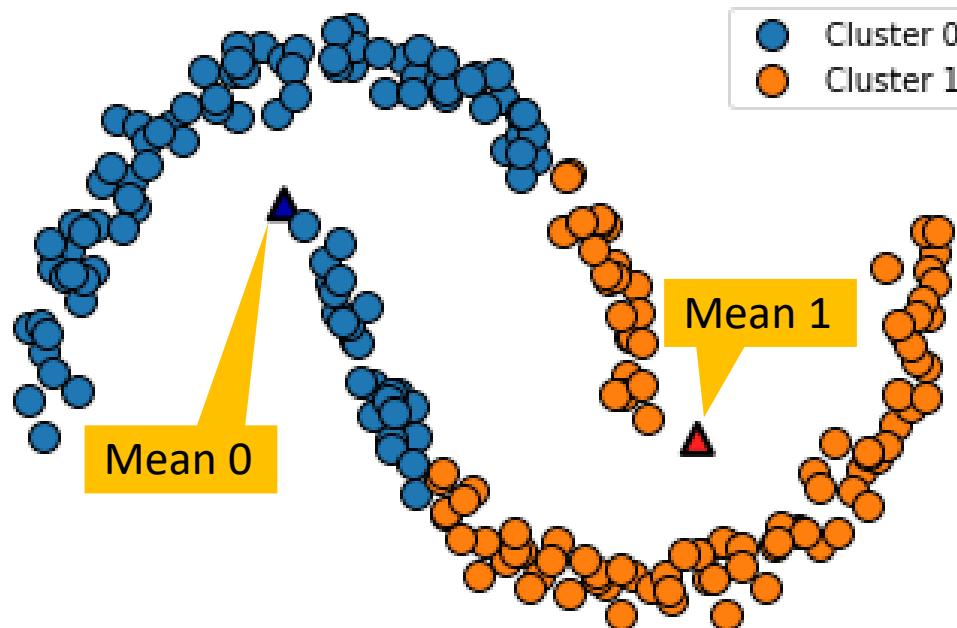
---

- **Weaknesses of K-Means**

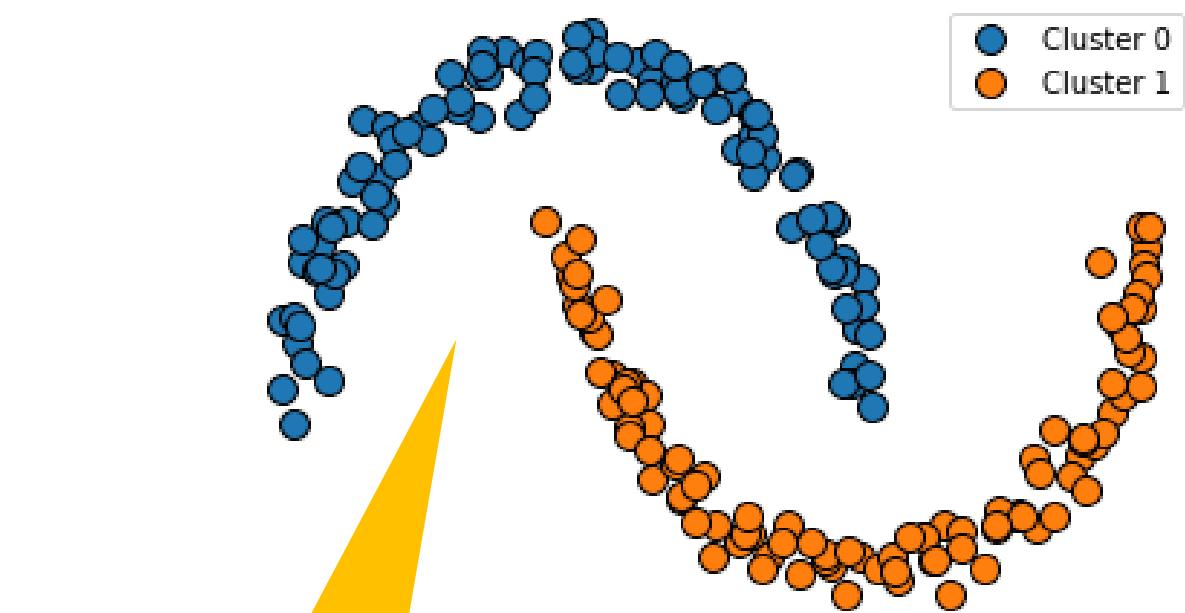
- Applicable only to points in a *continuous* n-dimensional space
  - We cannot calculate means on categorical values.
- Initialization matters.
  - Need to specify K, the number of clusters, in advance
    - In the literature, there are ways to automatically determine the best k
  - Different random initializations can create different final clusterings – K-means is only locally optimal
- Convergence
  - Stop condition can be ‘Relatively few points change the clusters’.
  - Often terminates at a *local* optimal.
- Sensitive to noisy data and outliers
- Not suitable to discover clusters with non-convex shapes

# K-means clustering

## K-means on an example of non-convex Shapes



K-means clustering result (K=2)

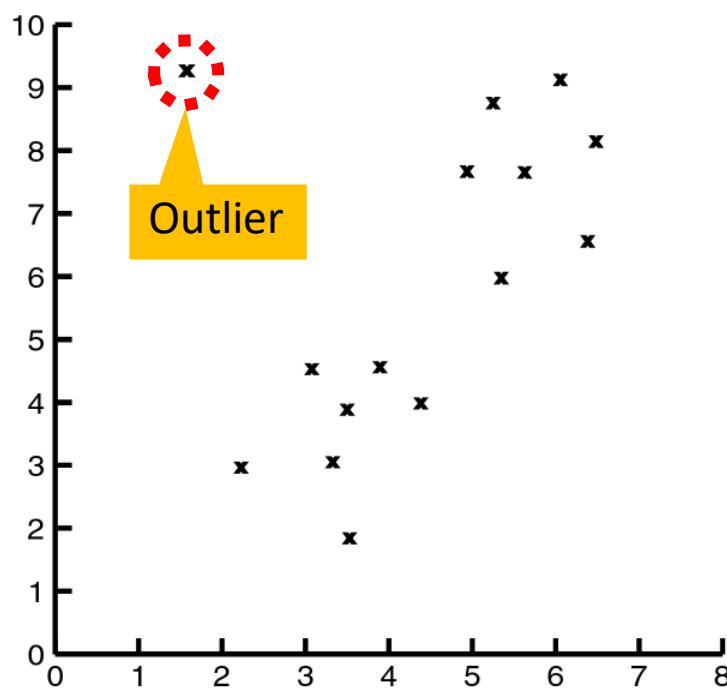


Desired clustering result

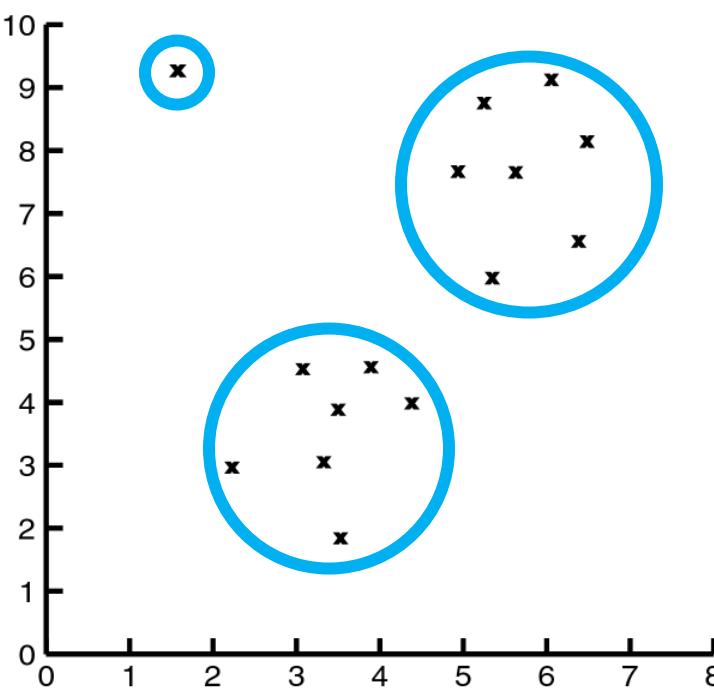
Density Based  
Spatial Clustering  
of Applications  
with Noise (DBSCAN)

# K-means clustering

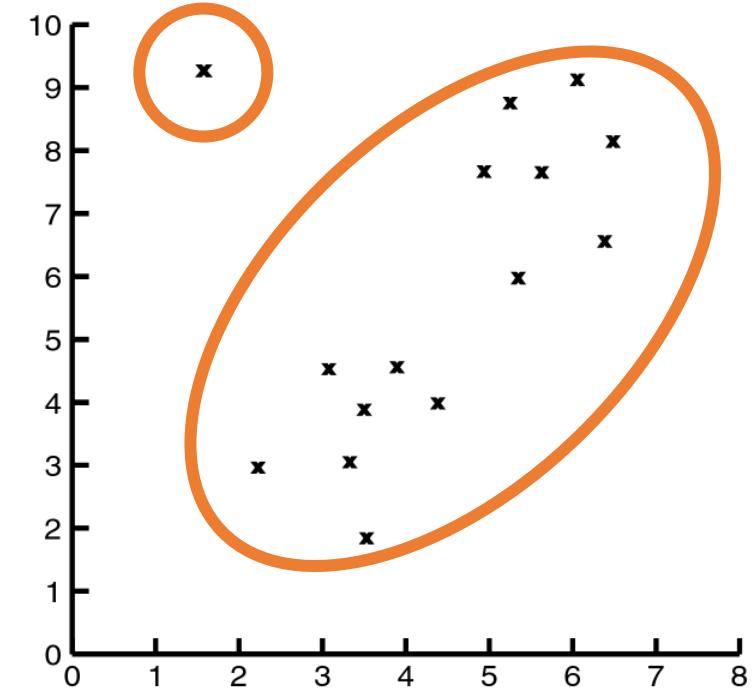
- Impact of Outliers on k-Means



Dataset with an outlier



K=3



K=2

# K-means clustering

---

- Let us look at examples in the notebook “K-Means clustering.ipynb”.

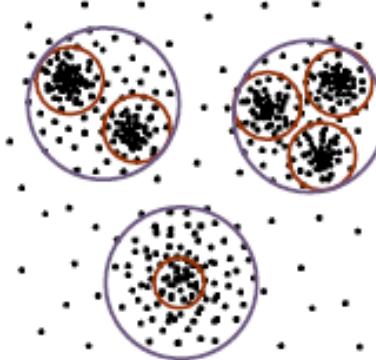
# Outline of this lecture

---

- Clustering as an example of unsupervised learning
- K-means clustering
- Hierarchical clustering
- DBSCAN clustering
- Evaluation of clustering models

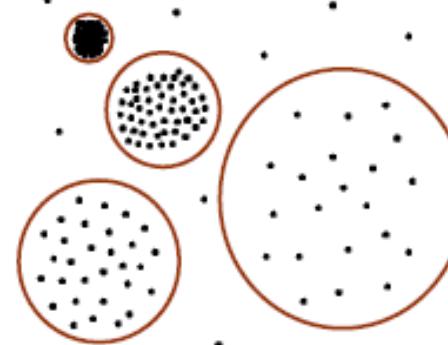
# Hierarchical clustering

- **Why Hierarchical Clustering?**
  - No need to specify the number of clusters beforehand



and/or

hierarchical  
cluster structure



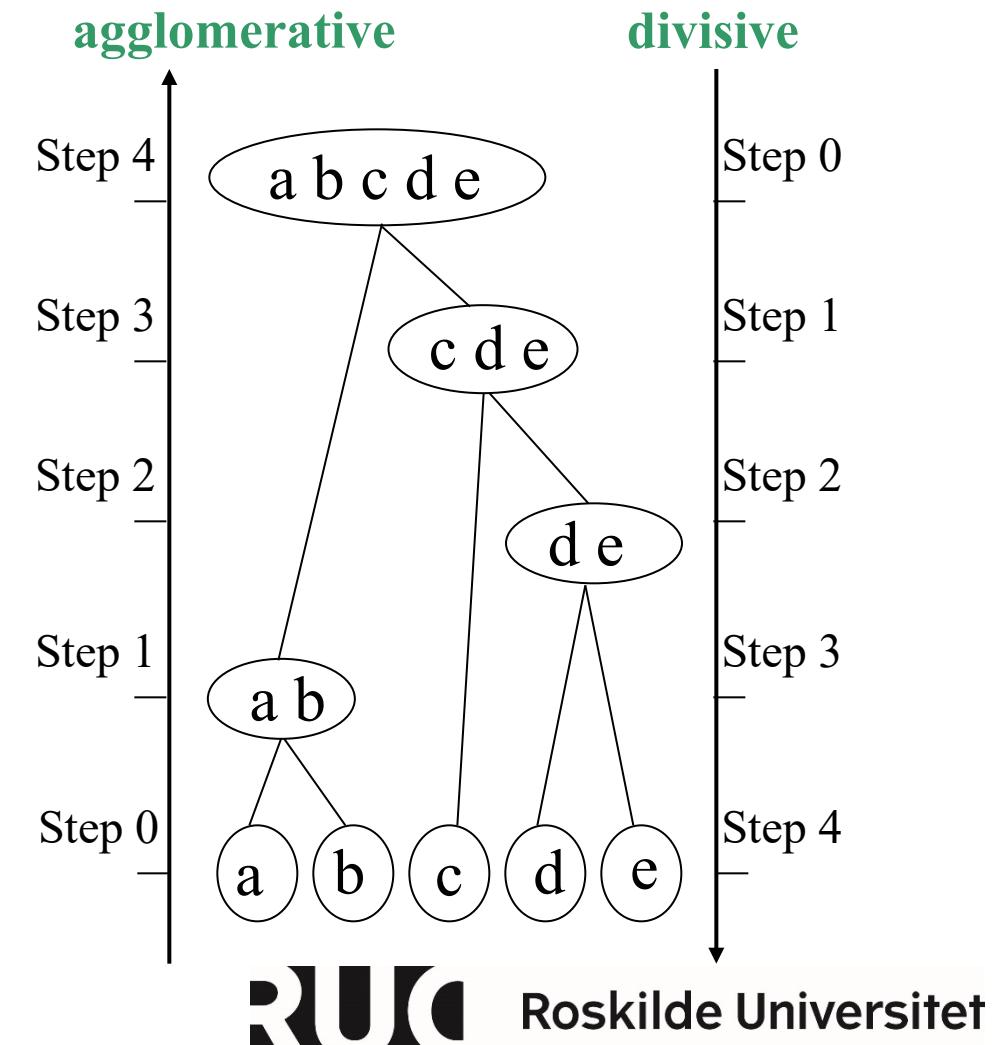
largely differing  
densities and sizes

- Hierarchical clustering can handle such situations.
  - Clusters are created in *levels*, creating sets of clusters at each level.

# Hierarchical clustering

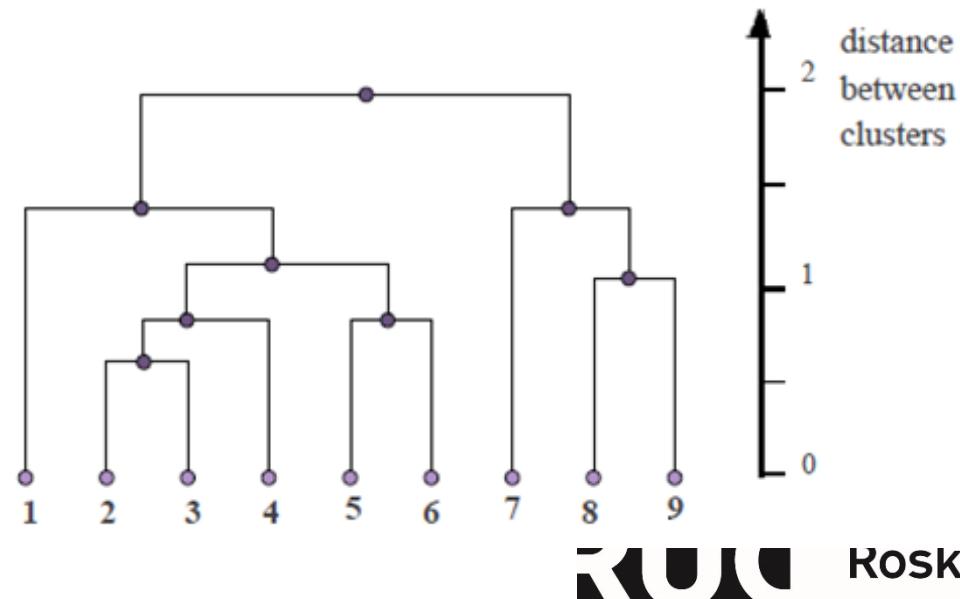
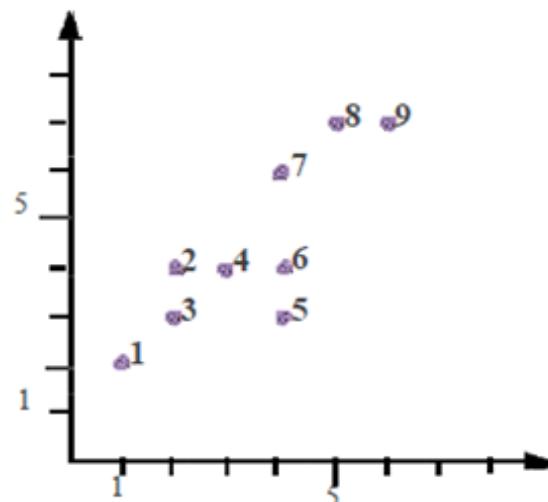
- **Hierarchical Clustering Approaches**

- Hierarchical Clustering needs a *distance metrics*, but do not need a fixed number of desired clusters in advance.
- Output: a hierarchy of potential clusterings
- **Agglomerative** clustering algorithms
  - Initially each item in its own cluster
  - Iteratively clusters are merged together
  - Bottom Up
- **Divisive** clustering algorithms
  - Initially all items in one cluster
  - Large clusters are successively divided
  - Top Down



# Hierarchical clustering

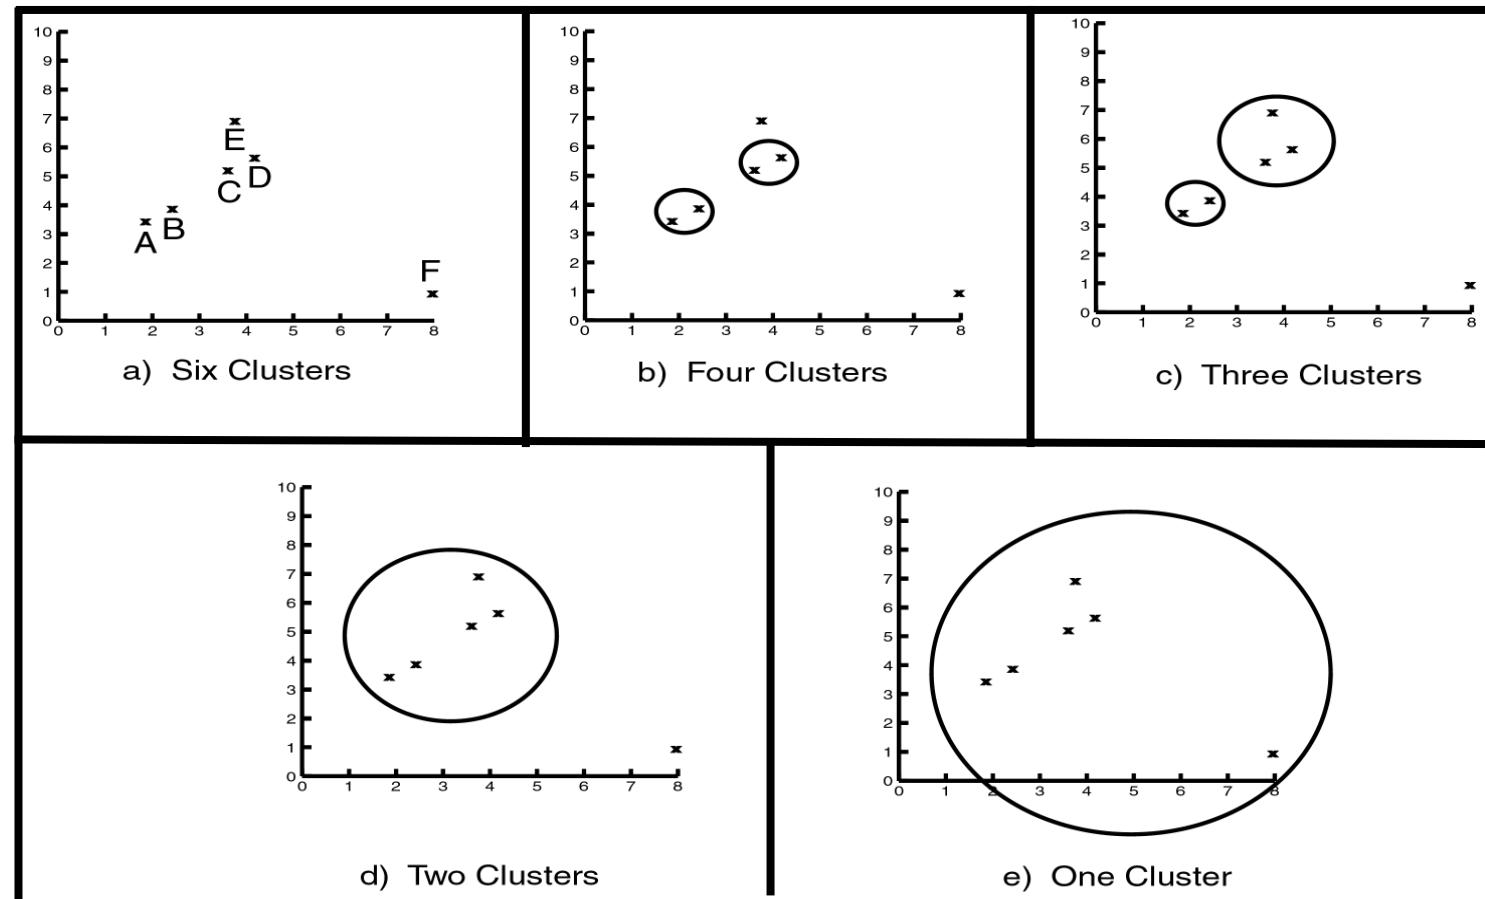
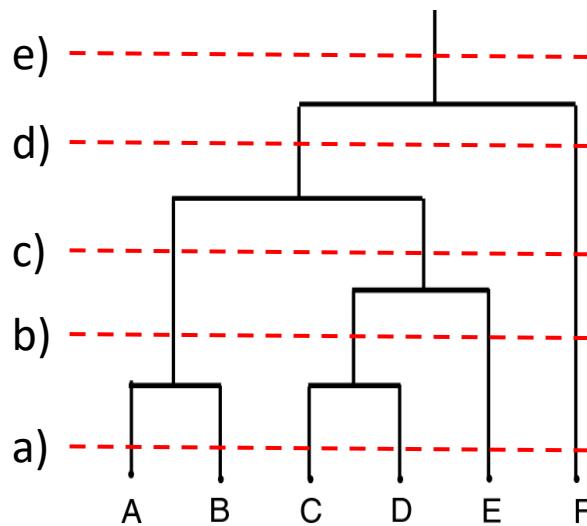
- **Dendrogram:** a tree data structure that illustrates hierarchical clustering techniques.
  - Each level shows clusters for that level.
    - Leaf: individual data points
    - Root: one cluster
    - A cluster at level  $i$  is the union of its child clusters at level  $i+1$ .
  - The height of an internal node represents the distance between its two child nodes.



# Levels of Clustering (Agglomerative)

## Levels of Clustering (Agglomerative)

- Each horizontal line in the Dendrogram, correspond to a particular set of clusters
  - where the vertical lines are intersected, the subtrees corresponds to the cluster
- The higher we cut the fewer clusters



# Hierarchical clustering

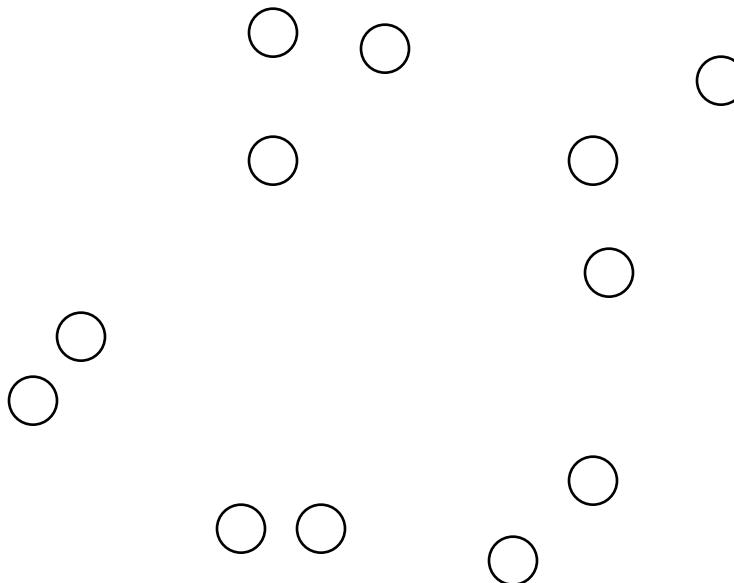
---

- **The Agglomerative Clustering Algorithm**
  - Most popular hierarchical clustering technique
  - **Basic algorithm:**
    1. Compute an *adjacency matrix*
    2. Let each data point be a cluster
    3. **Repeat**
      1. *Merge* two clusters if the distance is small enough
      2. *Update* the adjacency matrix
    4. **Until** only a single cluster remains
  - Key operation: *computing similarity of two clusters*
    - Different ways to define distance between clusters produce different clustering results.

# Hierarchical clustering

- **Agglomerative Clustering**

- Start with clusters of individual points and an adjacency matrix



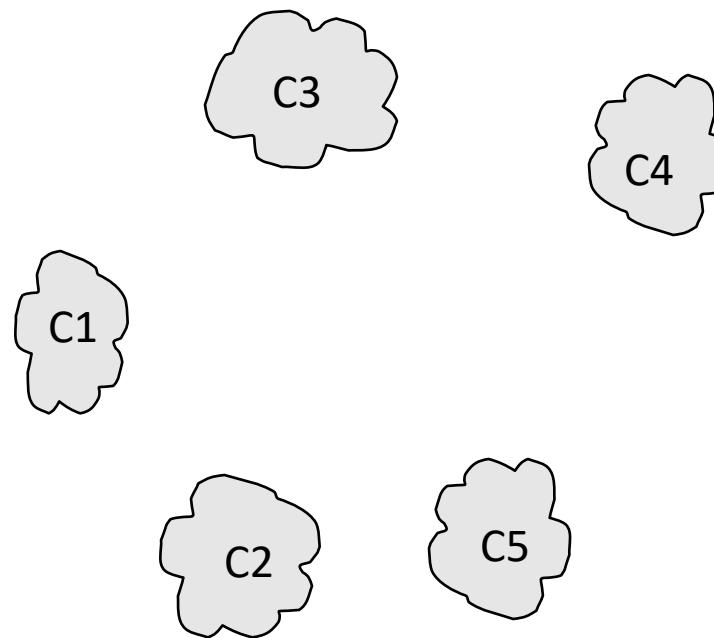
	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						

Adjacency Matrix

p1 p2 p3 p4 ... p9 p10 p11 p12

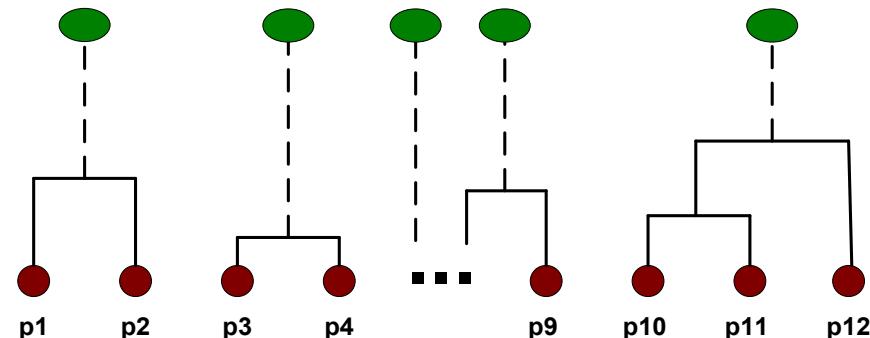
# Hierarchical clustering

- After some merging steps, we have some clusters



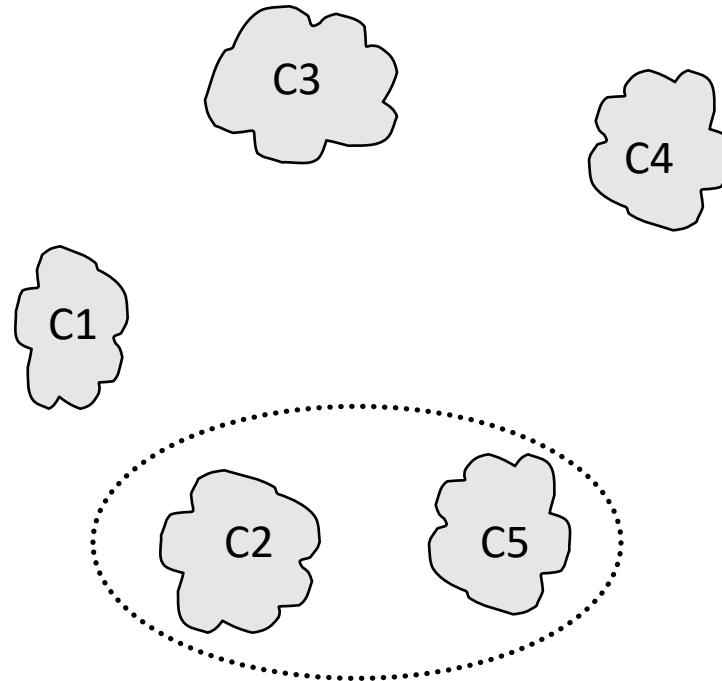
	c1	c2	c3	c4	c5
c1					

Adjacency Matrix



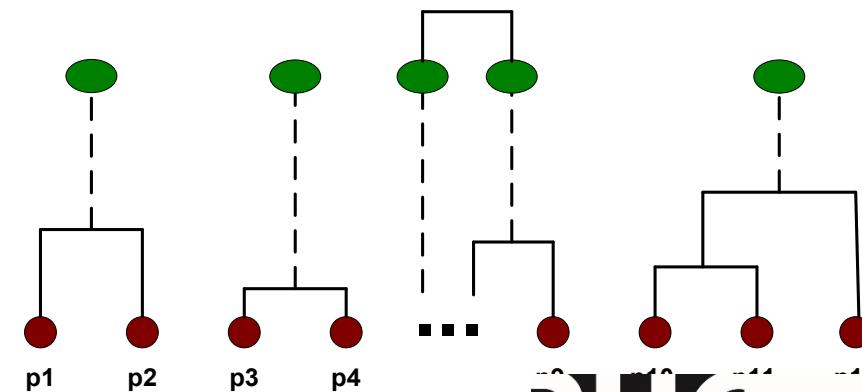
# Hierarchical clustering

- We want to merge two closest clusters (C2 and C5) and update the adjacency matrix.



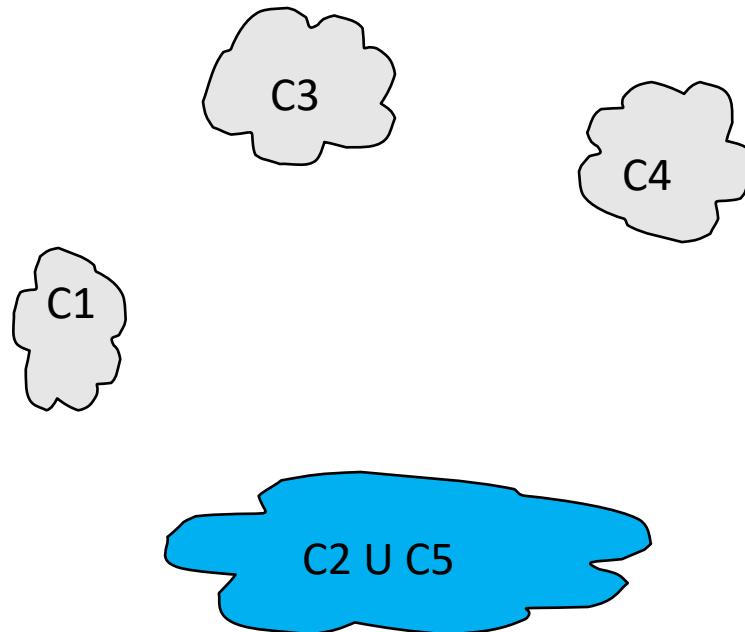
	c1	c2	c3	c4	c5
c1					

Adjacency Matrix



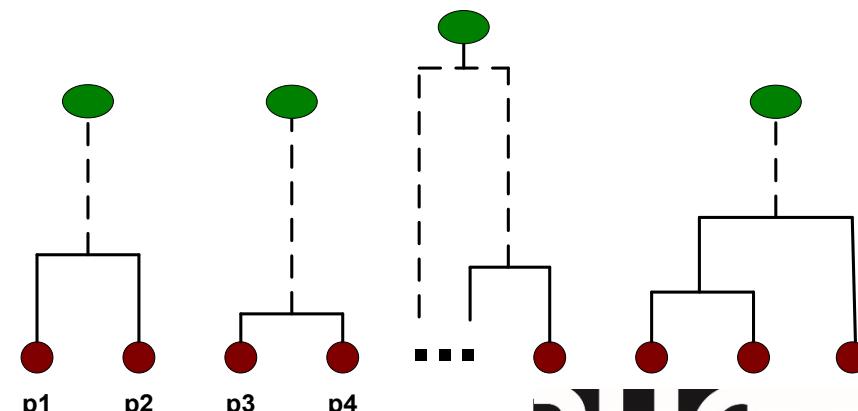
# Hierarchical clustering

- How to update the adjacency matrix?

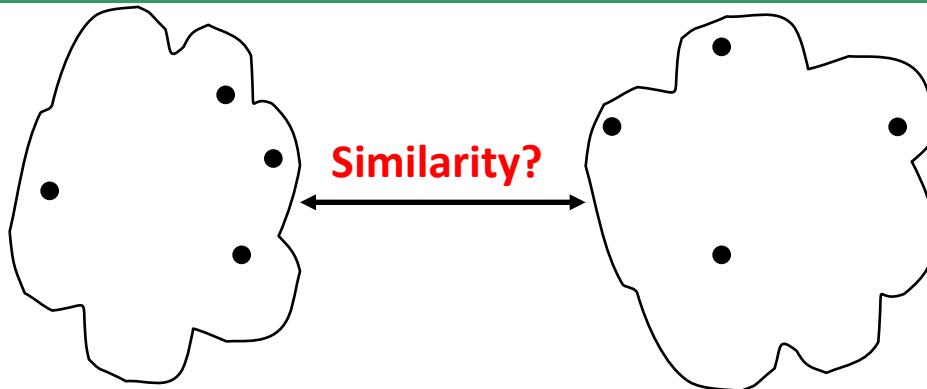


	C1	C5	C3	C4
C1		?		
C2 U C5	?	?	?	?
C3		?		
C4		?		

Proximity Matrix



# Hierarchical clustering



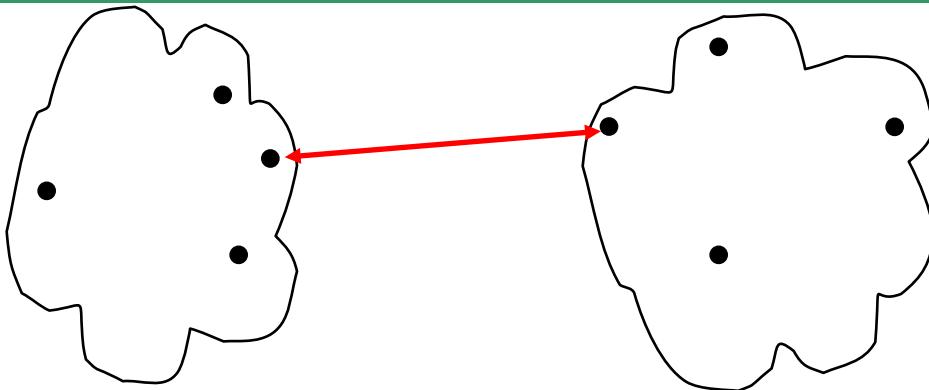
- **How to Define Inter-Cluster Similarity? (/linkage)?**

- Single (MIN)
- Complete (MAX)
- Average
- Centroid

	p1	p2	p3	p4	p5	...
p1						
.	.	.	.	.	.	.

- Adjacency Matrix

# Hierarchical clustering



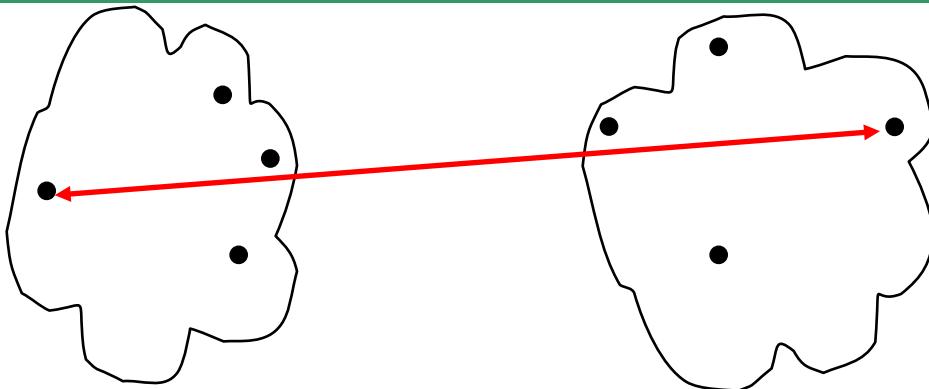
- **How to Define Inter-Cluster Similarity (/linkage)?**

- Single (MIN)
- Complete (MAX)
- Average
- Centroid

	p1	p2	p3	p4	p5	...
p1						
.	.	.	.	.	.	.

Adjacency Matrix

# Hierarchical clustering



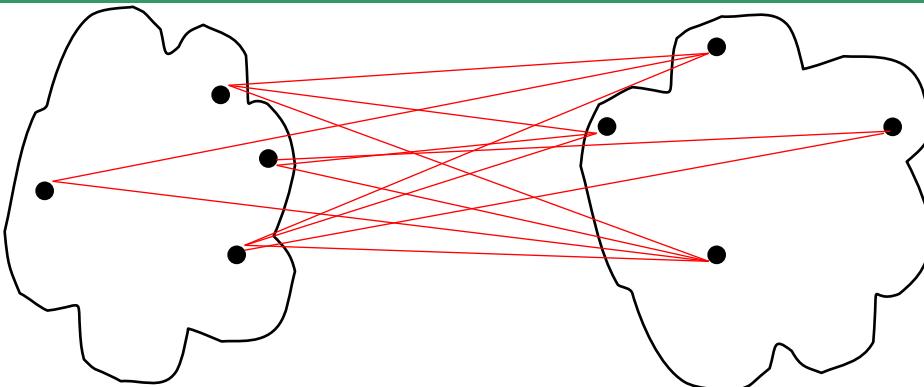
- **How to Define Inter-Cluster Similarity? (/linkage)?**

- Single (MIN)
- Complete (MAX)
- Average
- Centroid

	p1	p2	p3	p4	p5	...
p1						
.	.	.	.	.	.	.

Adjacency Matrix

# Hierarchical clustering



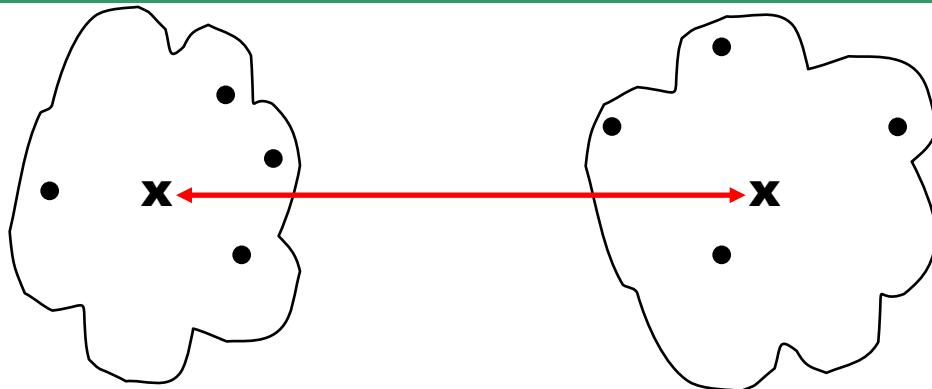
- **How to Define Inter-Cluster Similarity? (/linkage)?**

- Single (MIN)
- Complete (MAX)
- Average
- Centroid

	p1	p2	p3	p4	p5	...
p1						
.	.	.	.	.	.	.

Adjacency Matrix

# Hierarchical clustering



	p1	p2	p3	p4	p5	...
p1						
.	.	.	.	.	.	.

- **How to Define Inter-Cluster Similarity? (/linkage)?**

- Single (MIN)
- Complete (MAX)
- Average
- **Centroid**

- Adjacency Matrix

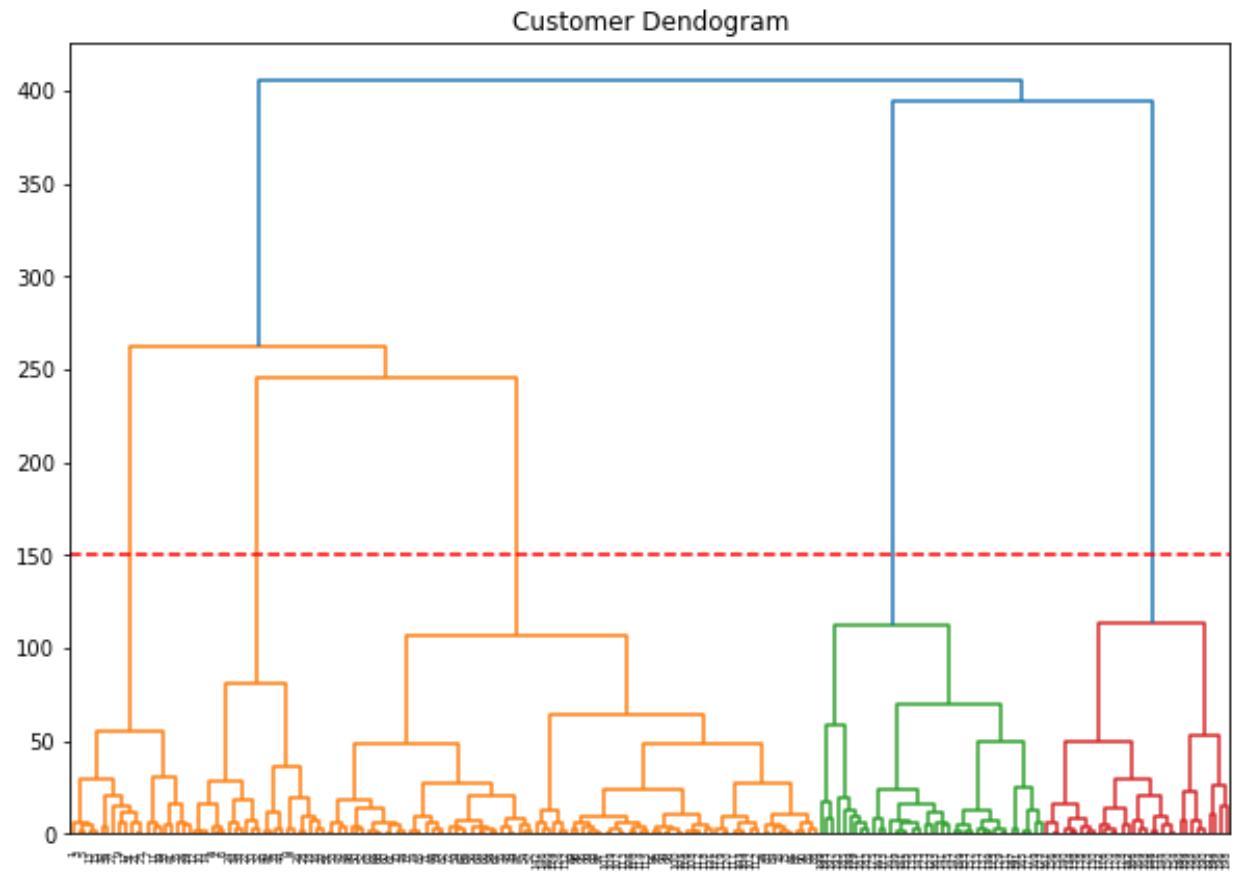
# Hierarchical clustering

---

- **Another Inter-Cluster Similarity: Ward's Method**
  - Similarity of two clusters measured as **increase** in **sum of squared error (SSE)** when they are merged
  - Less susceptible to noise and outliers
  - Biased towards globular clusters
  - Hierarchical “analogue” of K-means
  - Default in Scikit-Learns’s *AgglomerativeClustering* method:
    - <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.AgglomerativeClustering.html>

# Hierarchical clustering

- **Deciding the number of Clusters from a Dendrogram**
  - Locate the largest vertical difference between nodes
    - Avoid to merge very distant or dissimilar clusters
  - Draw a horizontal line through it.
    - If more options, choose the largest vertical difference again
  - Count the vertical lines it intersects
    - The *optimal* number of clusters.



# Hierarchical clustering

---

- Let us look at examples in the notebook “Hierarchical clustering.ipynb”.

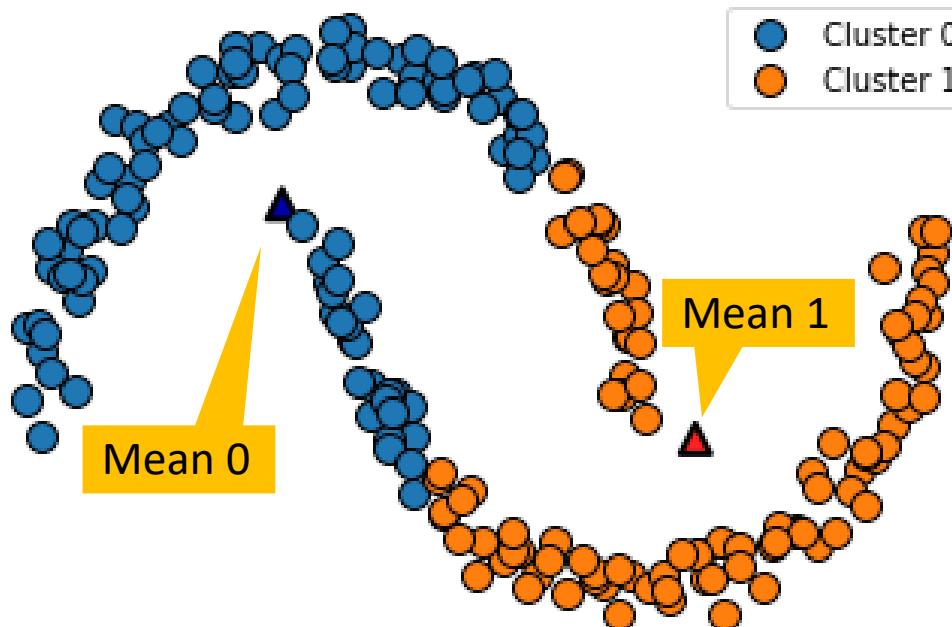
# Outline of this lecture

---

- Clustering as an example of unsupervised learning
- K-means clustering
- Hierarchical clustering
- DBSCAN clustering
- Evaluation of clustering models

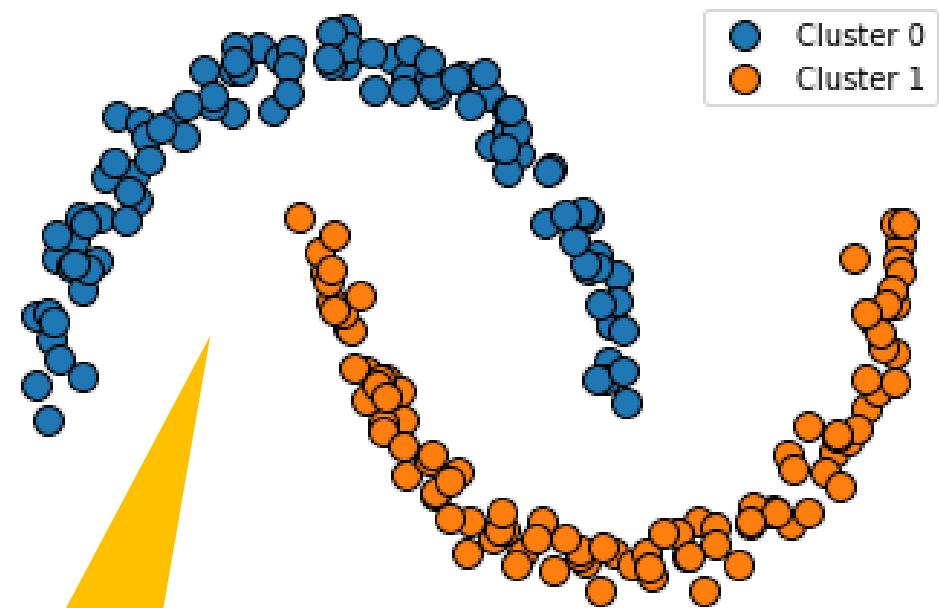
# DBSCAN clustering

Recall this example – how do we cluster the “right” way here



K-means clustering result (K=2)

Density Based  
Spatial Clustering  
of Applications  
with Noise (DBSCAN)



Desired clustering result

# DBSCAN clustering

---

- **DBSCAN: Density Based Spatial Clustering of Applications with Noise**
  - The algorithm's parameters (hyper-parameters):
    - **MinPts** – minimum number of points in a cluster
      - Size of a cluster (number of points)
      - **min\_samples** in `sklearn.cluster.DBSCAN`
    - **Eps** – for each point in a cluster there must be at least another one point in it less than this distance away.
      - Distance between points
      - **eps** in `sklearn.cluster.DBSCAN`

# DBSCAN clustering

- **Eps-neighborhood**

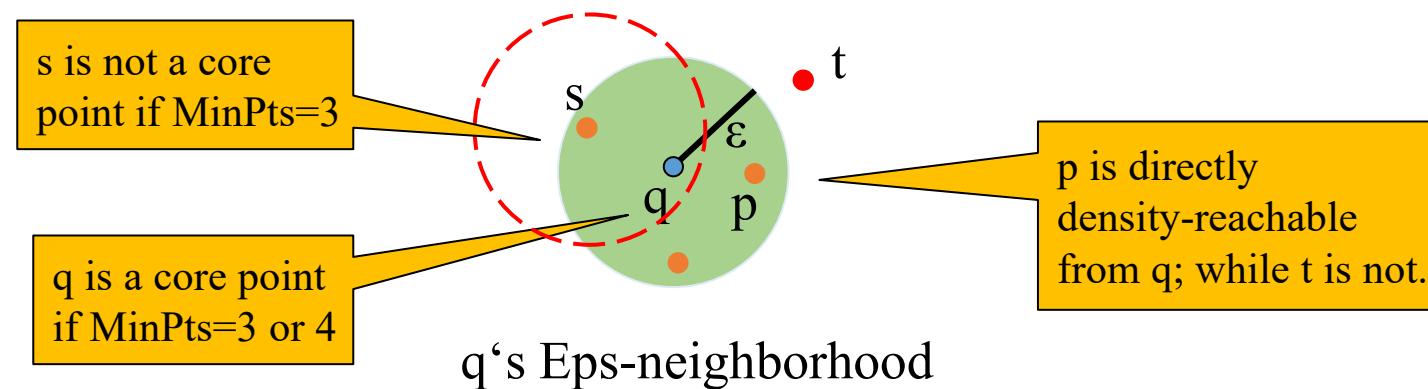
- Given a point, its Eps-neighborhood is all points within Eps distance of the given point.

- **Core point**

- Points whose Eps-neighborhood is dense enough (with at least MinPts points)

- **Directly density-reachable**

- A point  $p$  is directly density-reachable from another point  $q$  if the distance is small ( $\leq \text{Eps}$ ) and  $q$  is a core point.

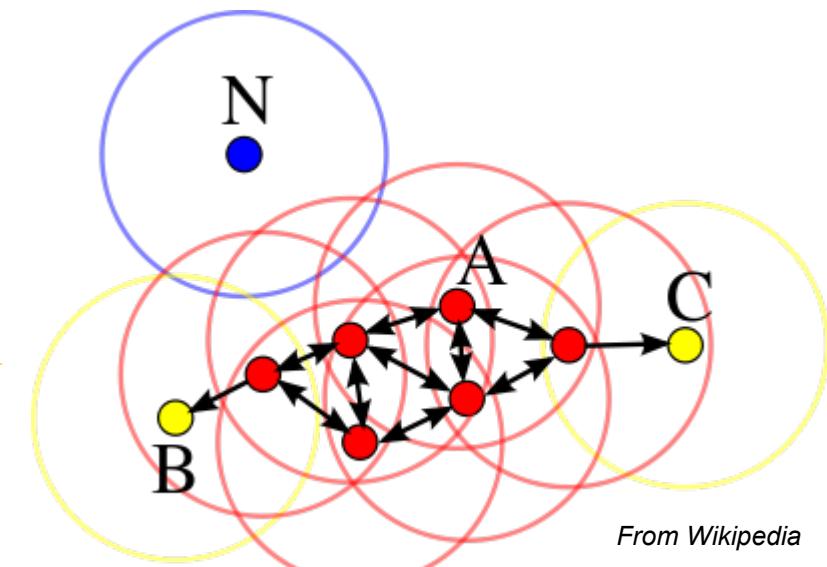


# DBSCAN clustering

- **Density-reachable:** A point  $p$  is density-reachable from another point  $q$  if there is a *path* from  $q$  to  $p$  and the path consists of only core points.
  - I.e., if there is a chain of points  $p_1=q, p_2, \dots, p_n=p$  such that  $p_{i+1}$  is directly density-reachable from  $p_i$ . More specifically,
    1.  $p_1, \dots, p_{n-1}$  are core points;
    2. the distance between each pair  $\leq \text{Eps}$ ;
    3.  $p$  may not be a core point.
  - Density-reachable is *not* symmetric.
    - A is not density-reachable from B or C as they are not core.

Assume MinPts=3.

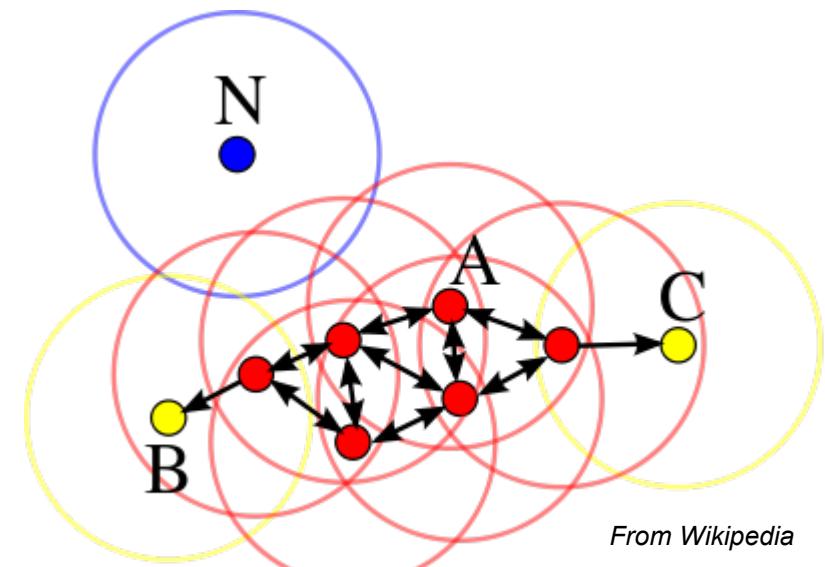
- Red points are core points.
- Points B and C are *density-reachable* from A.
- Point B is not density-reachable from C;  
and vice versa.



From Wikipedia

# DBSCAN clustering

- **Density-connected:** two points  $p$  and  $q$  are density-connected if there is a point  $o$  such that both  $p$  and  $q$  are density-reachable from  $o$ .
  - B and C are density-connected (via A).
  - Density-connected is symmetric.
- Clusters in DBSCAN
  - A cluster contains at least MinPts points
  - Density-connected points go to the same cluster
    - E.g., all red points plus B and C
- Outliers in DBSCAN
  - Those points that are not in any cluster
  - Outliers will note effect creation of clusters



From Wikipedia

# DBSCAN clustering

- The DBSCAN algorithm

```
DBSCAN(D, eps, MinPts)
    C = 0
    for each unvisited point P in dataset D
        mark P as visited
        NeighborPts = regionQuery(P, eps)
        if sizeof(NeighborPts) < MinPts
            mark P as NOISE ←
        else
            C = next cluster
            expandCluster(P, NeighborPts, C, eps, MinPts)

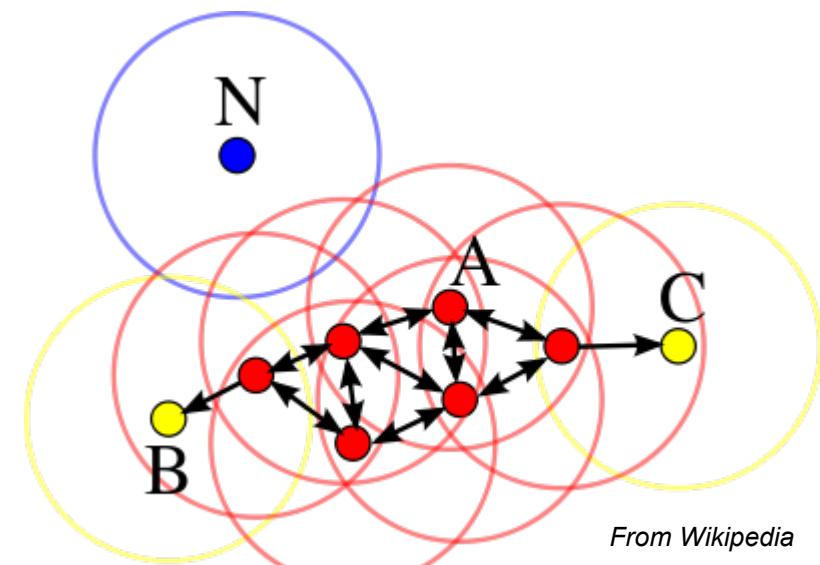
    → expandCluster(P, NeighborPts, C, eps, MinPts)
        add P to cluster C
        for each point P' in NeighborPts
            if P' is not visited
                mark P' as visited
                NeighborPts' = regionQuery(P', eps)
                if sizeof(NeighborPts') >= MinPts
                    NeighborPts = NeighborPts joined with NeighborPts'
            if P' is not yet member of any cluster
                add P' to cluster C

    → regionQuery(P, eps)
        return all points within P's eps-neighborhood
```

*From Wikipedia*

# DBSCAN clustering

- A cluster satisfies two properties:
  - All points within a cluster are mutually density-connected.
  - If a point  $p$  is density-connected to any point of a cluster,  $p$  belongs to the same cluster as well.
- In this example, point N is not included in any cluster.
  - It is a *noise point*, neither a core point nor density-reachable.



# DBSCAN clustering

---

- Let us look at examples in the notebook “DBSCAN clustering.ipynb”.

# Outline of this lecture

---

- Clustering as an example of unsupervised learning
- K-means clustering
- Hierarchical clustering
- DBSCAN clustering
- Evaluation of clustering models

# Evaluation of clustering models

---

- Quality: **What is a good clustering?**
  - A good clustering method will produce high quality clusters
    - high *intra-cluster* similarity: **cohesive** within clusters
    - low *inter-cluster* similarity: **distinctive** between clusters
  - The quality of a clustering method depends on
    - the similarity measure used by the method
    - its implementation (e.g., hyper-parameters), and
    - its ability to discover *some* or *all* of the hidden patterns

# Evaluation of clustering models

- **Rand Index (William M. Rand 1971) – measure difference between clusterings**
  - A set of points  $S = \{o_1, \dots, o_n\}$ . Two clusterings:  $X = \{X_1, \dots, X_r\}$  and  $Y = \{Y_1, \dots, Y_s\}$ 
    - $a$ : #pairs of elements in  $S$  that are in the **same**  $X_i$  and in the **same**  $Y_j$
    - $b$ : #pairs of elements in  $S$  that are in **different**  $X_i$ s and in **different**  $Y_j$ s
    - $c$ : #pairs of elements in  $S$  that are in the **same**  $X_i$  but in **different**  $Y_j$ s
    - $d$ : #pairs of elements in  $S$  that are in **different**  $X_i$ s but in the **same**  $Y_j$
  - Rand Index  $R = \frac{a+b}{a+b+c+d} = \frac{a+b}{\binom{n}{2}}$ , where  $\binom{n}{2} = \frac{n(n-1)}{2}$  (binomial coefficient)
    - A value between 0 and 1.
    - 0: the two clusterings do not agree on any pair of points.
    - 1: the two clusterings are exactly the same.
  - Example
    - Dataset: {A, B, C, D, E}
    - Method 1 Clusters: {{A, B, C}, {D, E}} , Method 2 Clusters: {{A, B}, {C, D}, {E}}
    - $a=1$ : {A, B};  $b=5$ : {A, D}, {A, E}, {B, D}, {B, E}, {C, E};  $a+b+c+d=\binom{5}{2}=10$
    - $R = (1+5)/10 = 0.6$

# Evaluation of clustering models

- **Adjusted Rand Index – measure difference between clusterings**

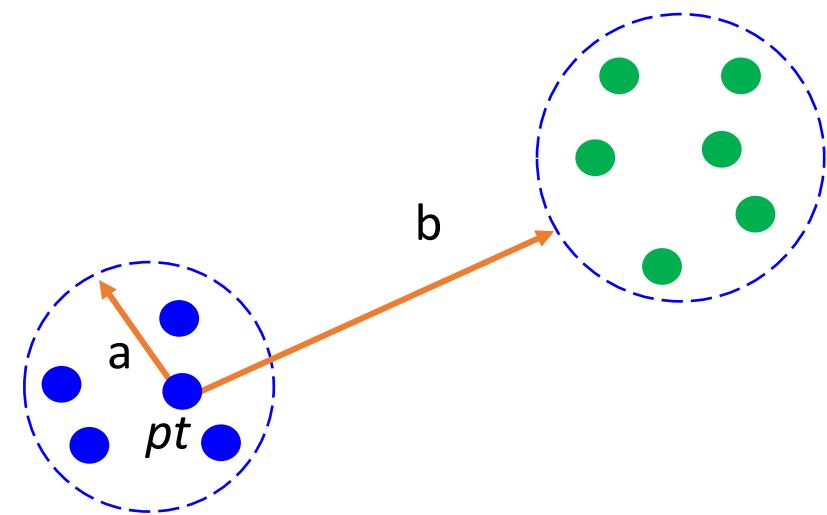
- A set of points  $S = \{o_1, \dots, o_n\}$ . Two clusterings:  $X = \{X_1, \dots, X_r\}$  and  $Y = \{Y_1, \dots, Y_s\}$
- **The contingency table:**  $n_{ij} = |X_i \cap Y_j|$ 
  - Each entry denotes the number of objects in common between  $X_i$  and  $Y_j$

$X \setminus Y$	$Y_1$	$Y_2$	$\cdots$	$Y_s$	sums
$X_1$	$n_{11}$	$n_{12}$	$\cdots$	$n_{1s}$	$a_1$
$X_2$	$n_{21}$	$n_{22}$	$\cdots$	$n_{2s}$	$a_2$
$\vdots$	$\vdots$	$\vdots$	$\ddots$	$\vdots$	$\vdots$
$X_r$	$n_{r1}$	$n_{r2}$	$\cdots$	$n_{rs}$	$a_r$
sums	$b_1$	$b_2$	$\cdots$	$b_s$	

- **Adjusted Rand Index:**  $ARI = \frac{\sum_{ij} \binom{n_{ij}}{2} - [\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2}] / \binom{n}{2}}{\frac{1}{2} [\sum_i \binom{a_i}{2} + \sum_j \binom{b_j}{2}] - [\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2}] / \binom{n}{2}}$

# Evaluation of clustering models

- **Silhouette Score – measure the quality of a clusterings**
  - Silhouette score for one point  $pt$ 
    - $s(pt) = (b - a) / \max(a, b)$
    - a: the **average** distance between  $pt$  and all others in the same cluster (**cohesive**)
    - b: the **smallest average** distance between  $pt$  and all points in any other cluster (**distinctive**)
  - Silhouette score for a clustering result  $X$ 
    - $s(X) = (\bar{b} - \bar{a}) / \max(\bar{a}, \bar{b})$ 
      - $\bar{a}, \bar{b}$ : Average a and b for all points in the dataset
    - 1: Clusters are well apart from each other and clearly distinguished.
    - 0: Clusters are indifferent. The distance between them is insignificant.
    - -1: Clusters are assigned in the wrong way.



# Evaluation of clustering models

---

- Evaluation of Clustering in Scikit-Learn
  - If you have clustering groundtruth
    - Compare the clustering result with the groundtruth by measuring a score
      - Adjusted Rand Index (**ARI**): `adjusted_rand_score(groundtruth, clustering_result)`
  - Otherwise
    - `silhouette_score(X, clustering_results)` computes the compactness of a cluster
  - All scores are in `sklearn.metrics.cluster`

# Evaluation of clustering models

---

- Let us look at examples in the notebook “Evaluation of clustering models.ipynb”.

## Exercises

- Do the exercise in the notebook “Exercises in Clustering.ipynb”